Making Predictions

MARKETING ANALYTICS: PREDICTING CUSTOMER CHURN IN PYTHON

Mark Peterson
Director of Data Science, Infoblox

datacamp

(Supervised) Machine Learning Primer

- Goal: Predict whether or not a customer will churn
- Target Variable: `'Churn'`
- Supervised Machine Learning
- Learn from historical (training) data to make new predictions


scikit
learn

Model selection

- Which model to use?
- ... it depends!
- In this course: Experiment with several models
- To learn about their inner workings: Check out other DataCamp courses

Model selection

- Logistic regression: Good baseline
  - Offers simplicity and interpretability
  - Cannot capture more complex relationships
- Random forests
- Support vector machines

Training your model

```
from sklearn.svm import SVC
svc = SVC()
svc.fit(telco[features], telco['target'])
```

```
SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.0,
  decision_function_shape='ovr', degree=3, gamma='scale', kernel='rbf',
  max_iter=-1, probability=False, random_state=None, shrinking=True,
  tol=0.001, verbose=False)
```

Making a prediction

```
prediction = svc.predict(new_customer)


print(prediction)
```

```
[0]
```

Let's practice!

MARKETING ANALYTICS: PREDICTING CUSTOMER CHURN IN PYTHON

Evaluating Model Performance

MARKETING ANALYTICS: PREDICTING CUSTOMER CHURN IN PYTHON

Mark Peterson
Director of Data Science, Infoblox

datacamp

Accuracy

- One possible metric: Accuracy
  - Total Number of Correct Predictions / Total Number of Data Points
- What data to use?
  - Training data not representative of new data

Training and Test Sets

- Fit your classifier to the training set
- Make predictions using the test set

Training and Test Sets using scikit-learn

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(telco['data'], telco['target'],
                                                    test_size=0.2, random_state = 42)
from sklearn.svm import SVC

svc = SVC()

svc.fit(X_train, y_train)

svc.predict(X_test)
```

Computing Accuracy

```
svc.score(X_test, y_test)
```

```
0.857
```

- 85.7% accuracy: Quite good for a first try!

Improving your model

- Overfitting: Model fits the training data too closely
- Underfitting: Does not capture trends in the training data
- Need to find the right balance between overfitting and underfitting

Let's practice!

MARKETING ANALYTICS: PREDICTING CUSTOMER CHURN IN PYTHON

Model Metrics

MARKETING ANALYTICS: PREDICTING CUSTOMER CHURN IN PYTHON

Mark Peterson
Director of Data Science, Infoblox

datacamp

Imbalanced classes

```
telco['Churn'].value_counts()
```

```
no     2850
yes     483
Name: Churn, dtype: int64
```

- Accuracy not a very useful metric

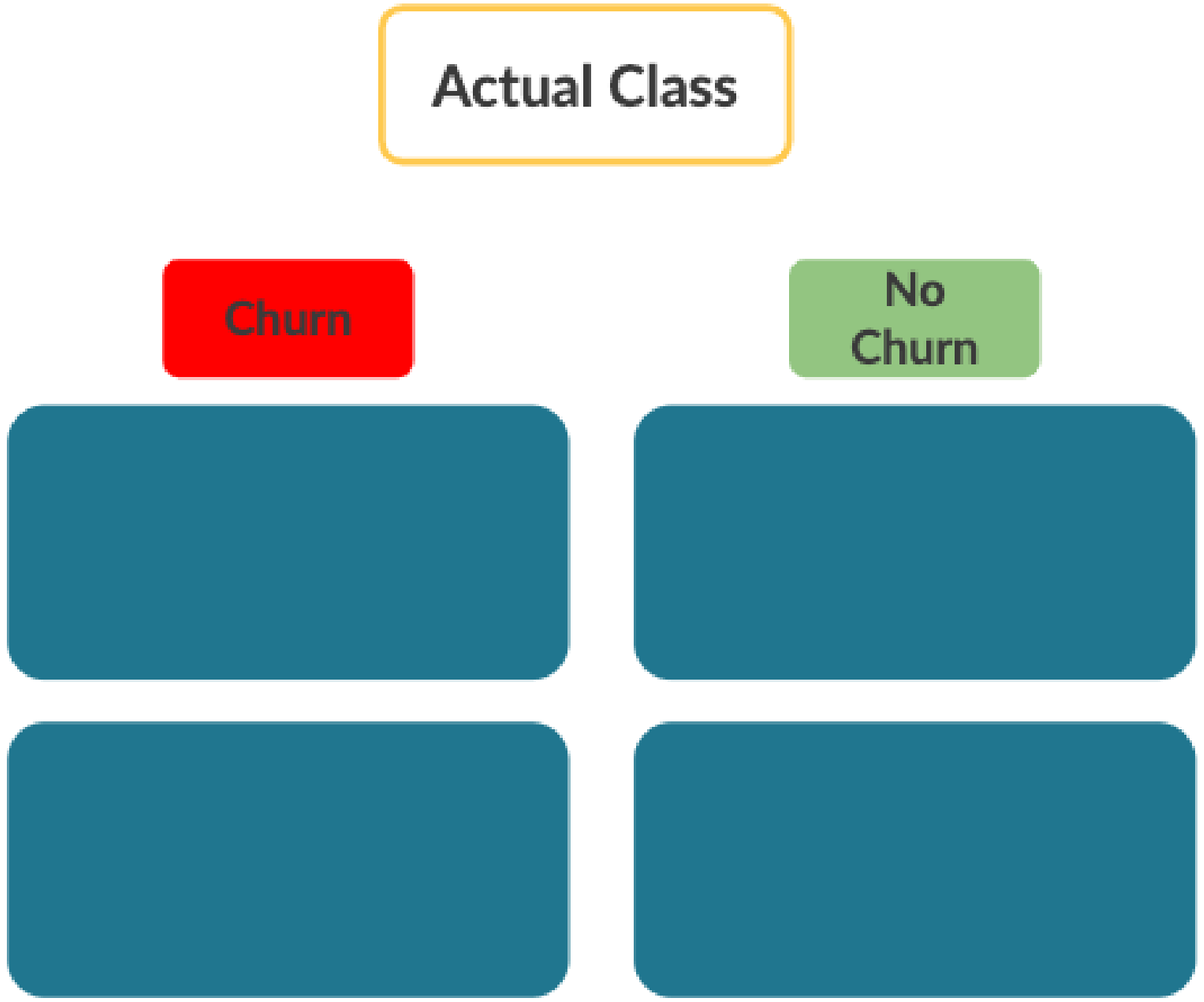
Actual Class
Churn
No Churn

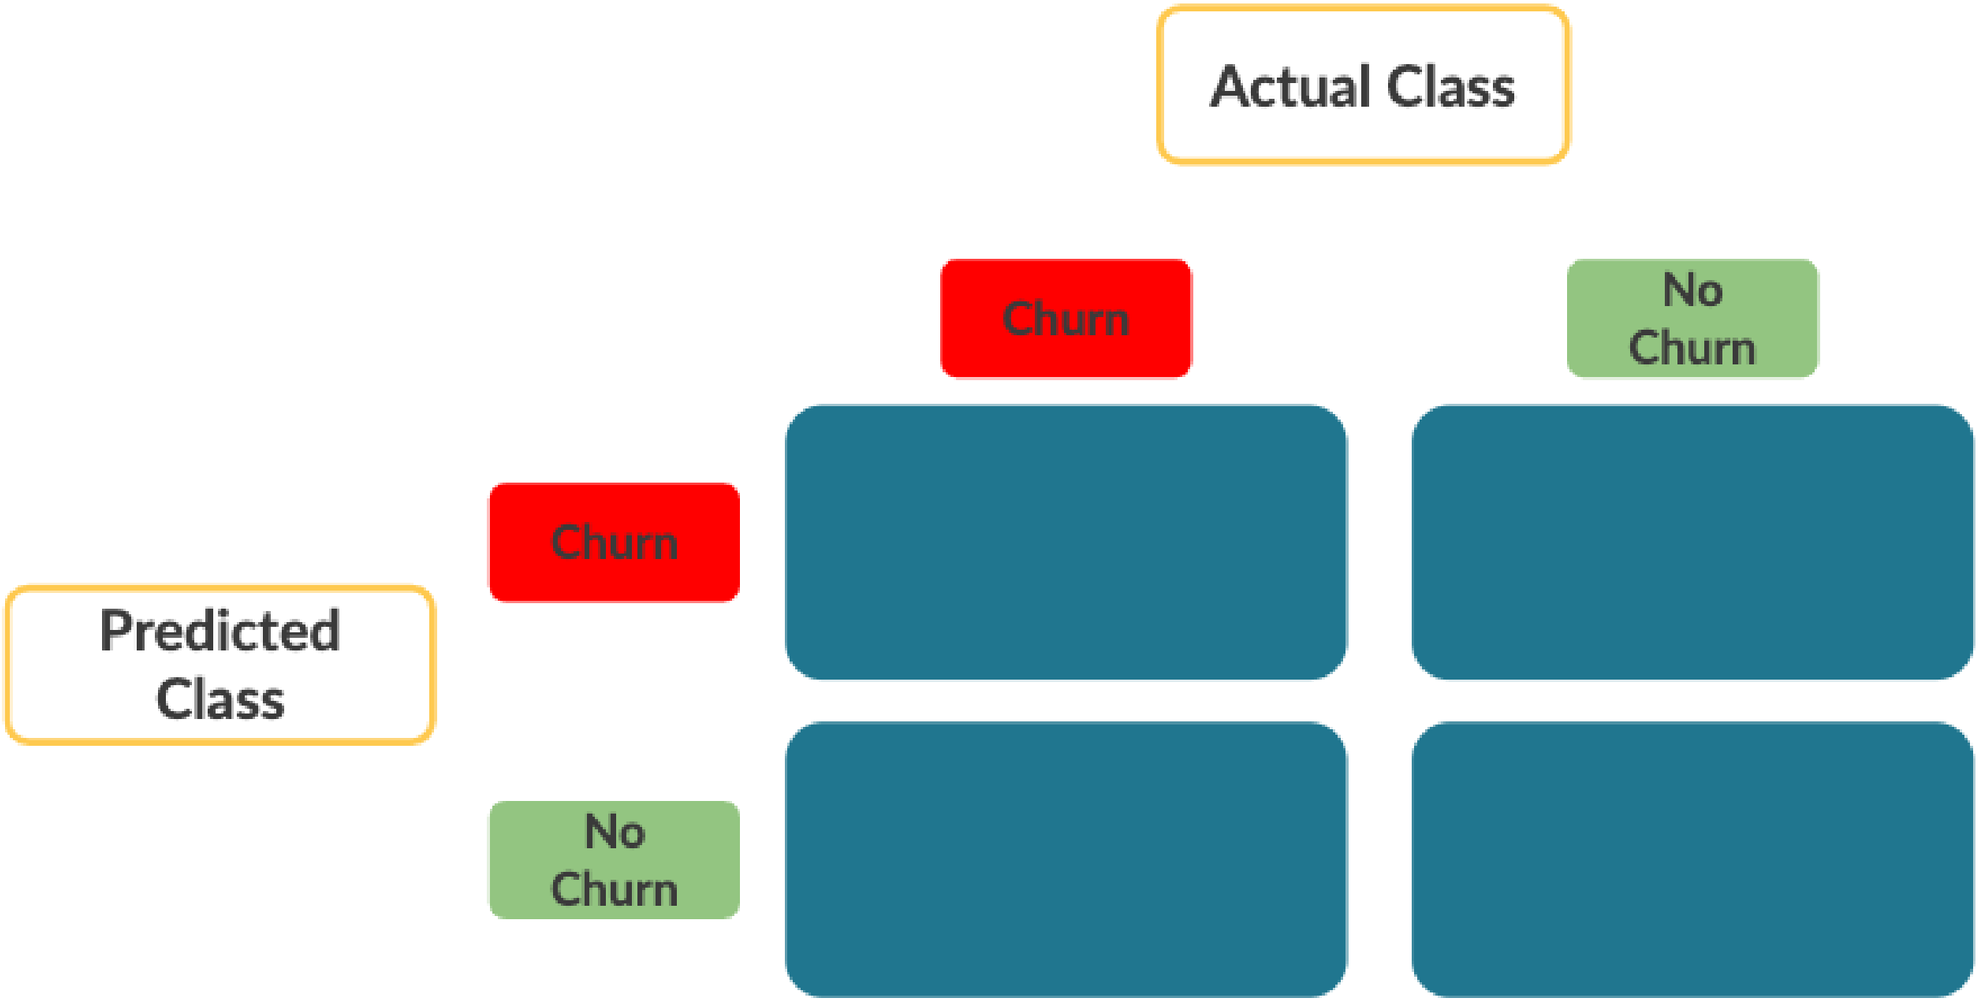
Actual Class
Churn
No Churn
Predicted Class
Churn
No Churn

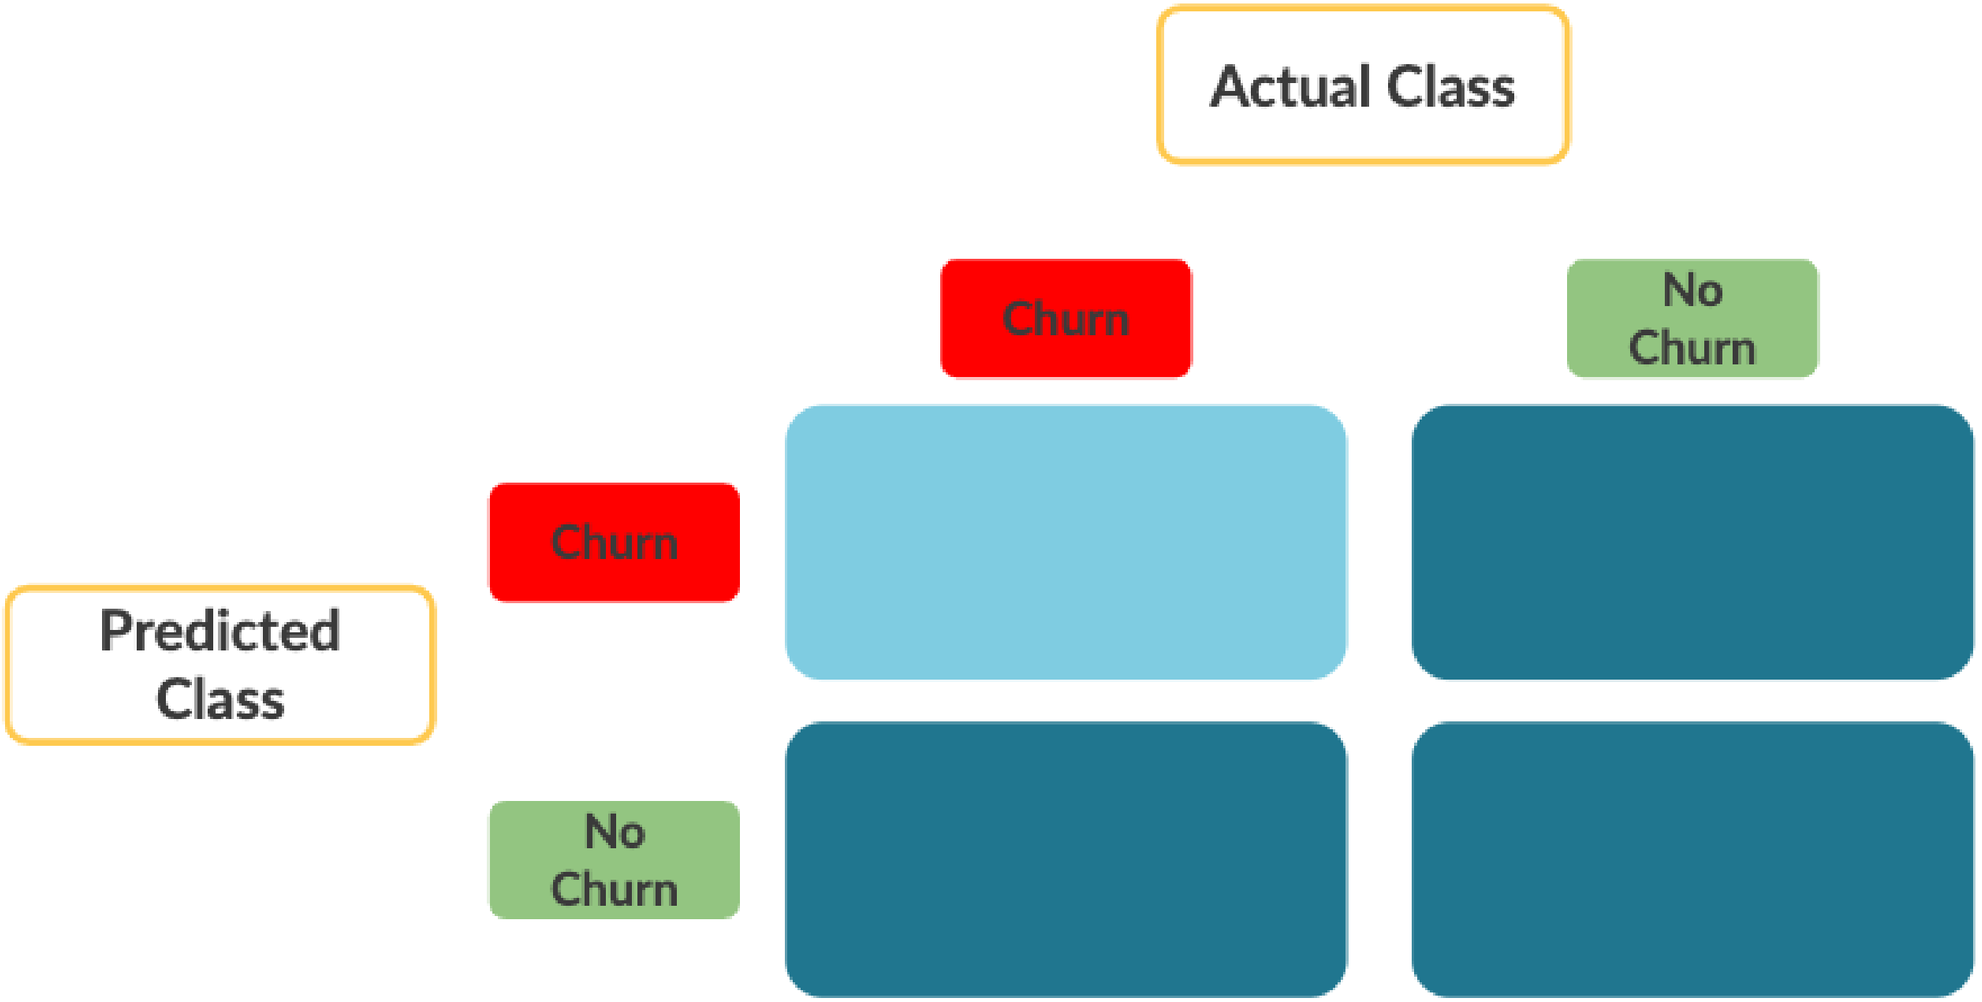
Actual Class
Churn
No Churn
Predicted Class
Churn
No Churn

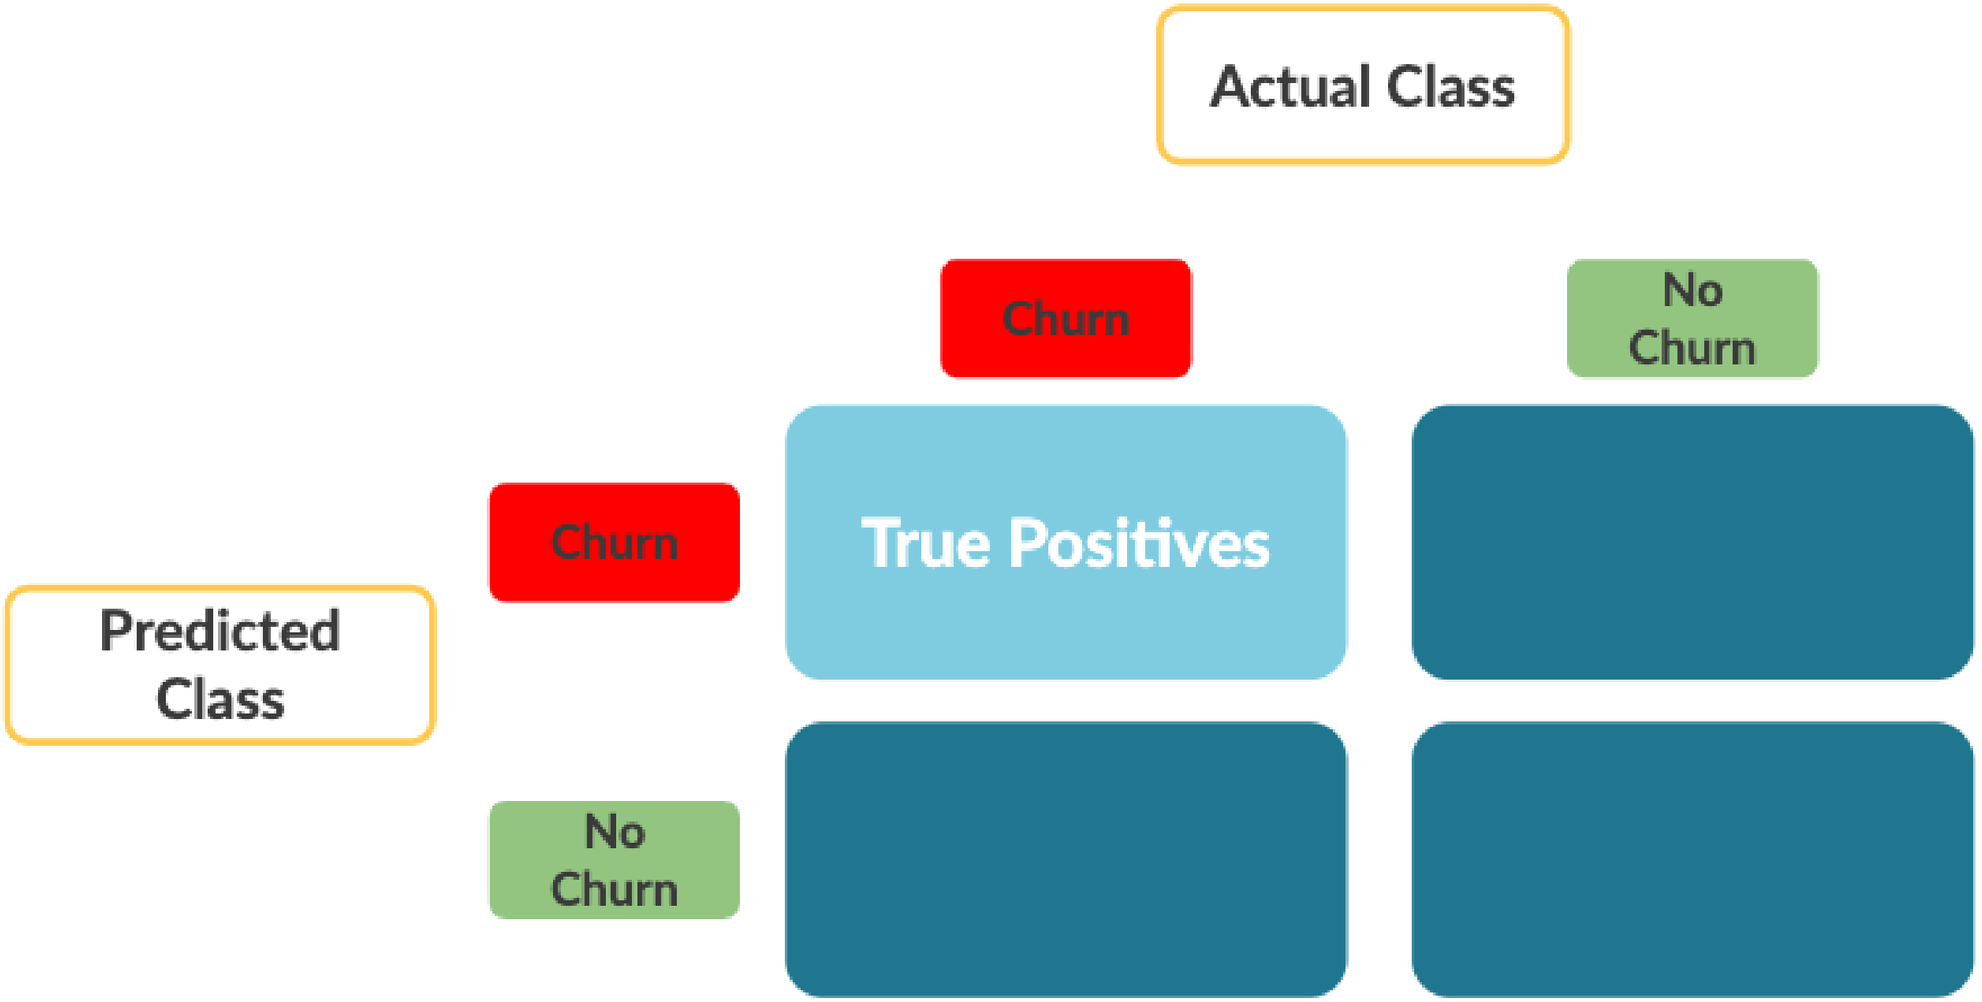
Actual Class
Churn
No Churn
Predicted Class
Churn
True Positives
No Churn

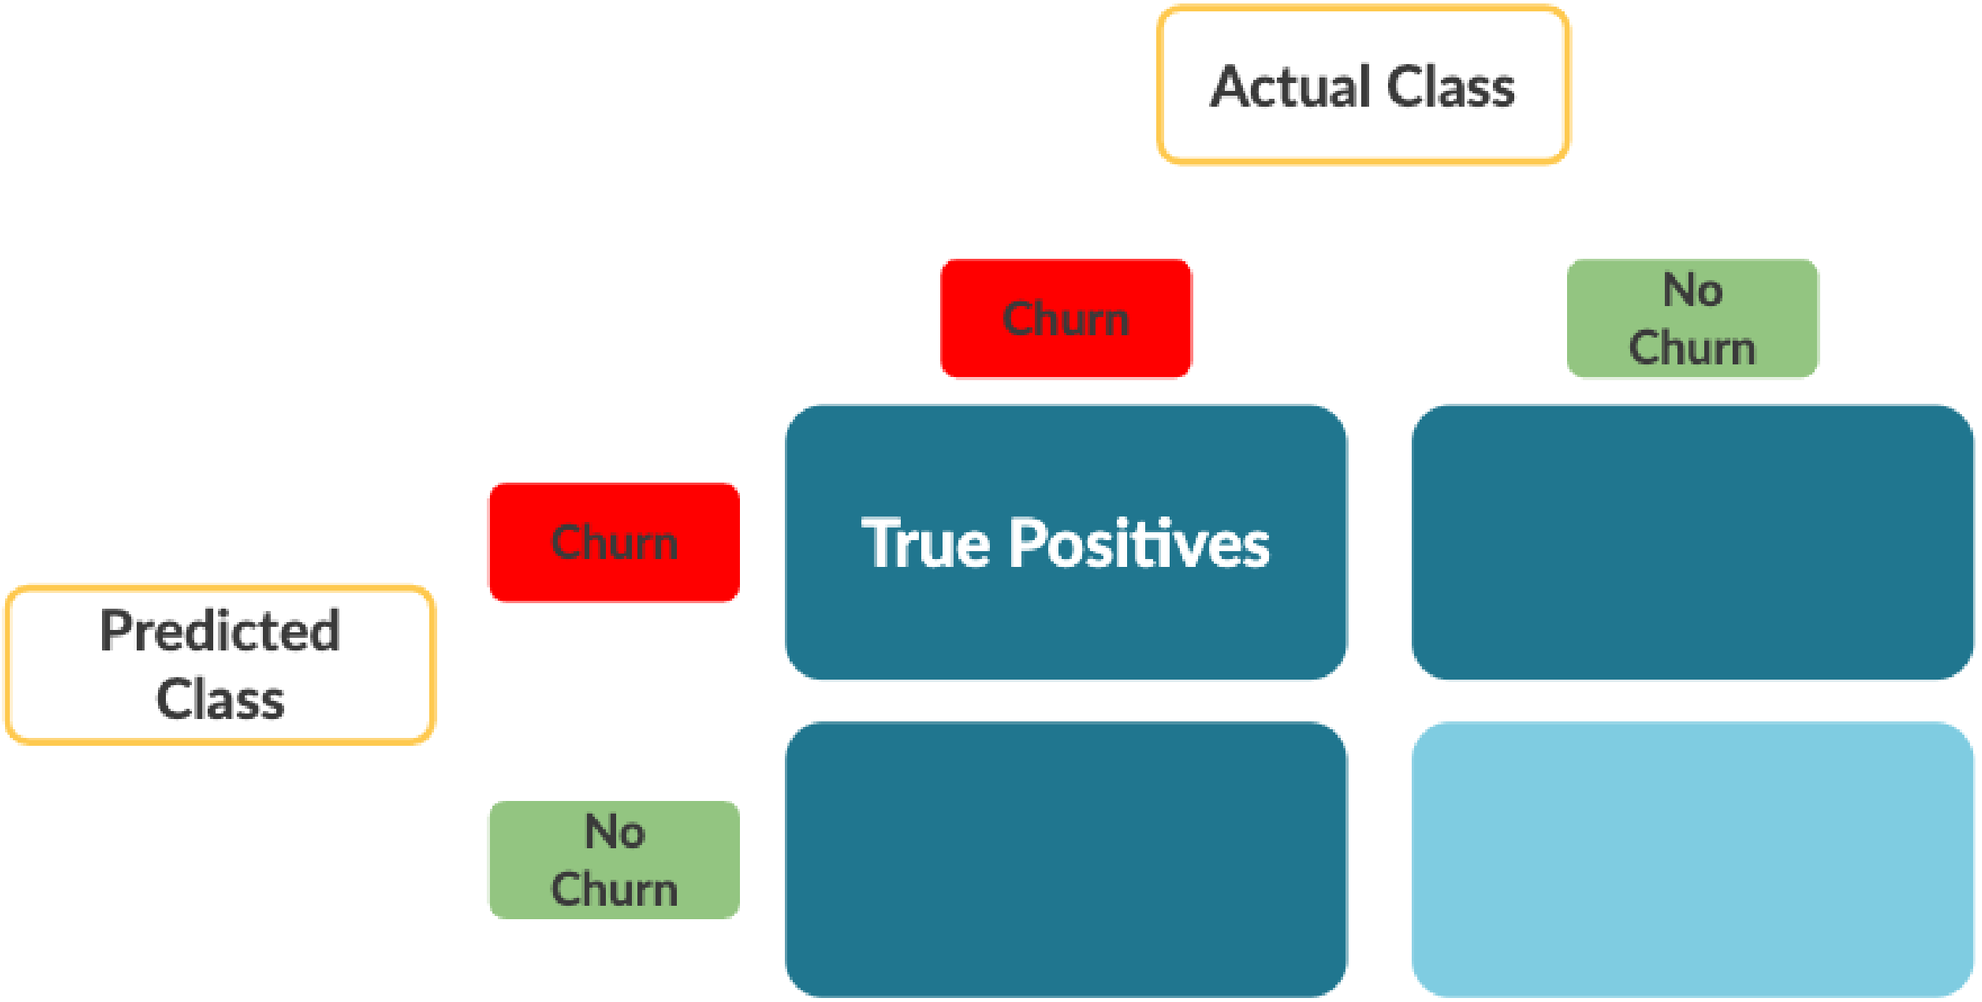
Actual Class
Churn
No Churn
Predicted Class
Churn
True Positives
No Churn

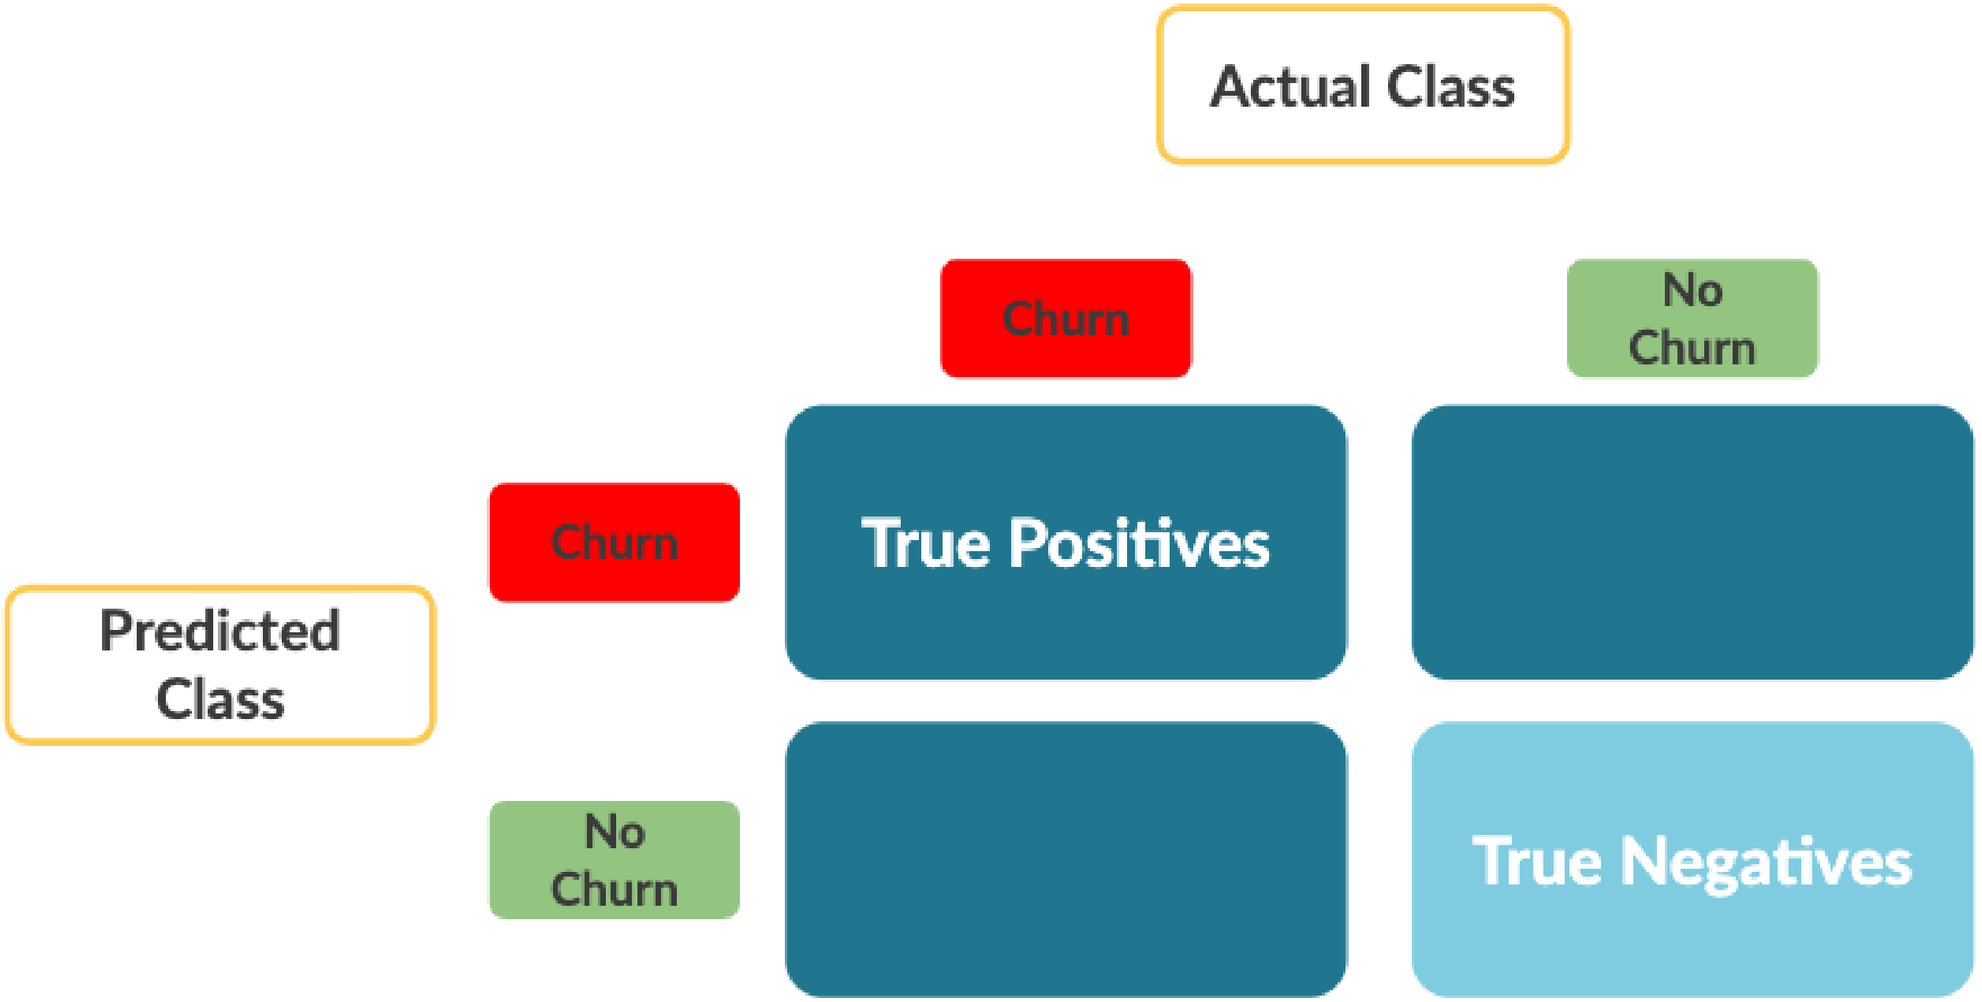
Actual Class
Churn
No Churn
Predicted Class
Churn
True Positives
No Churn
True Negatives

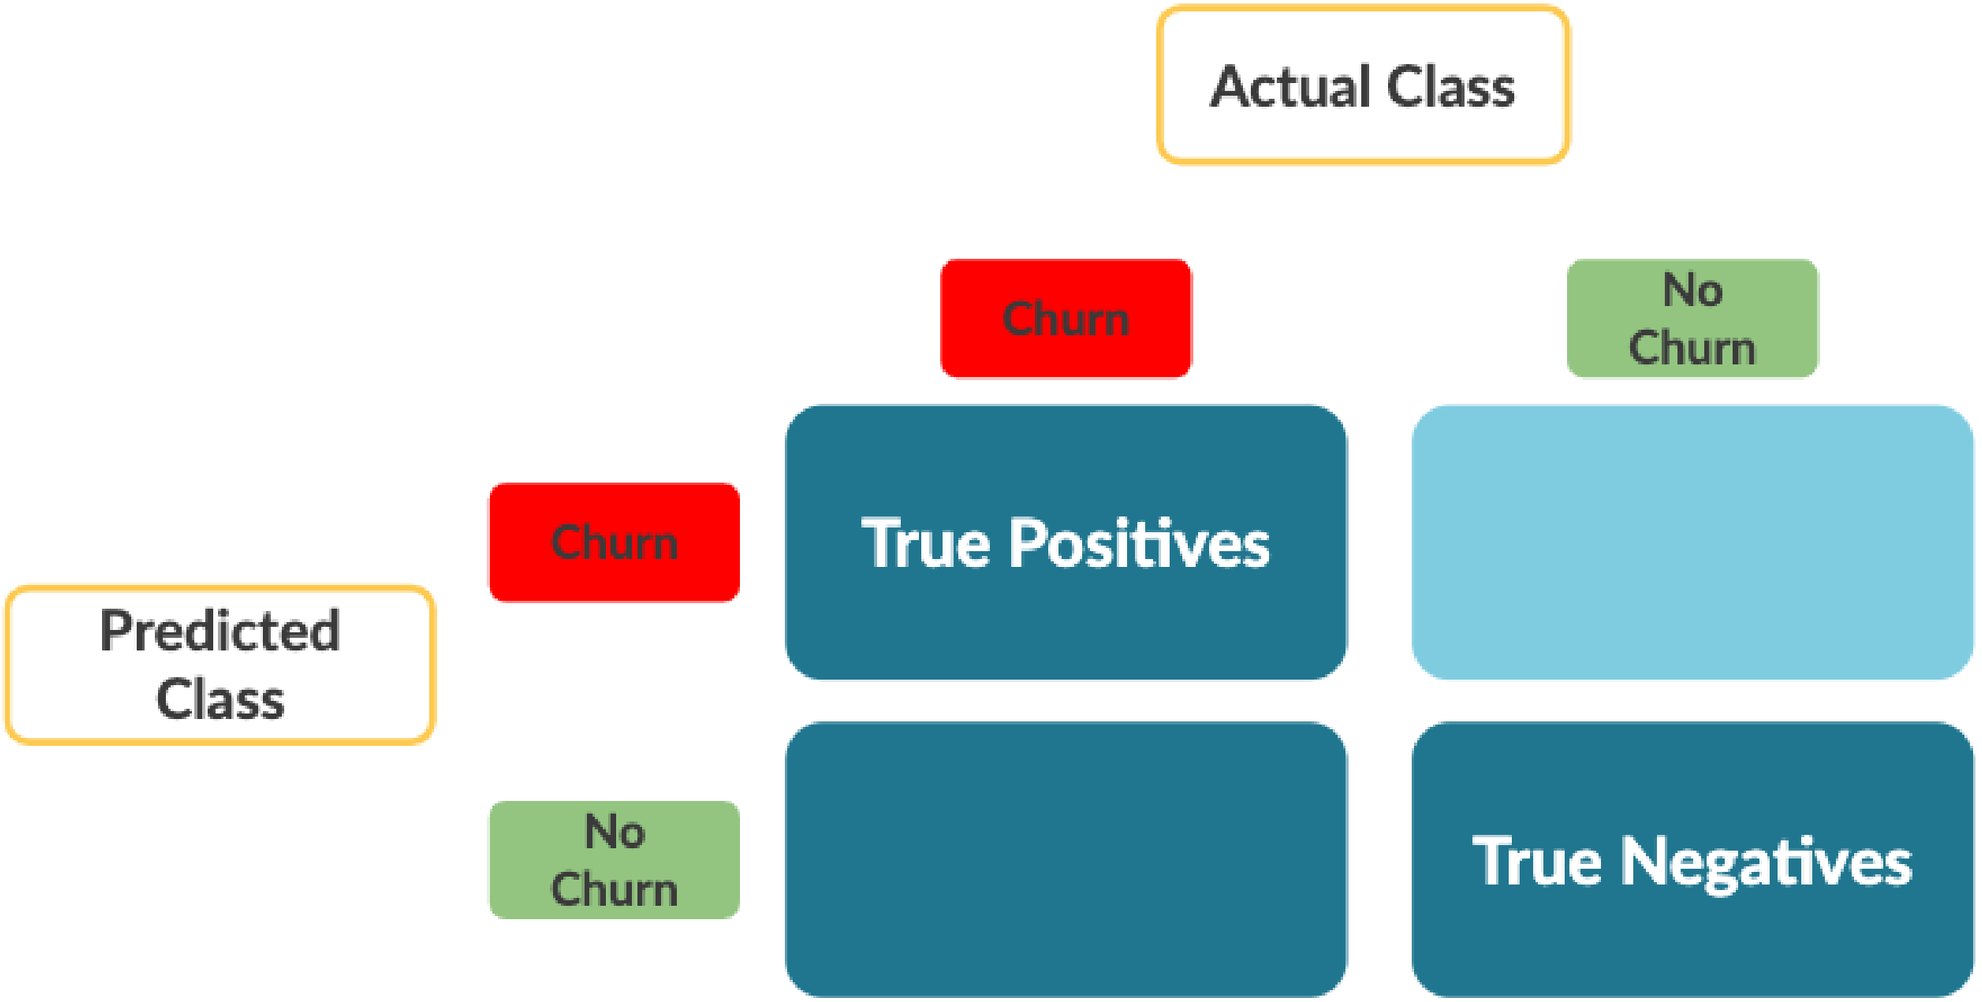
Actual Class
Churn
No Churn
Predicted Class
Churn
True Positives
No Churn
True Negatives

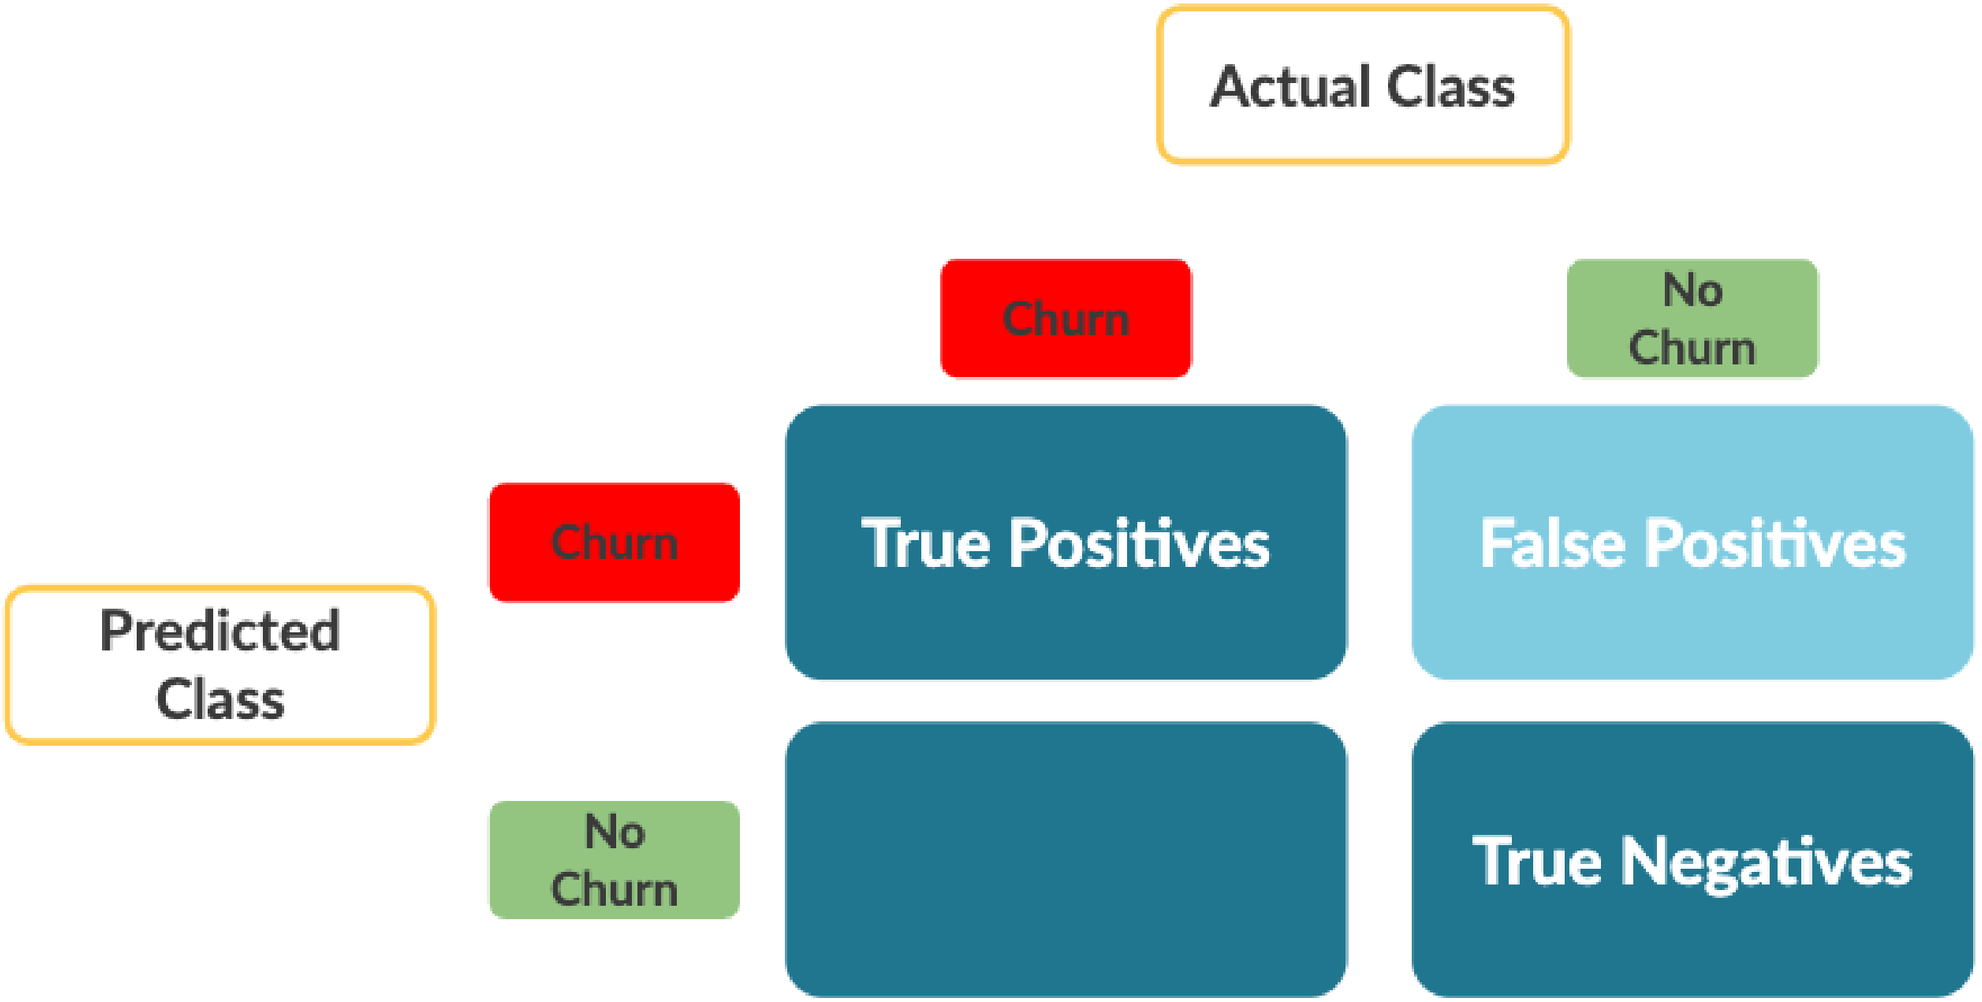
Actual Class
Churn
No Churn
Predicted Class
Churn
True Positives
False Positives
No Churn
True Negatives

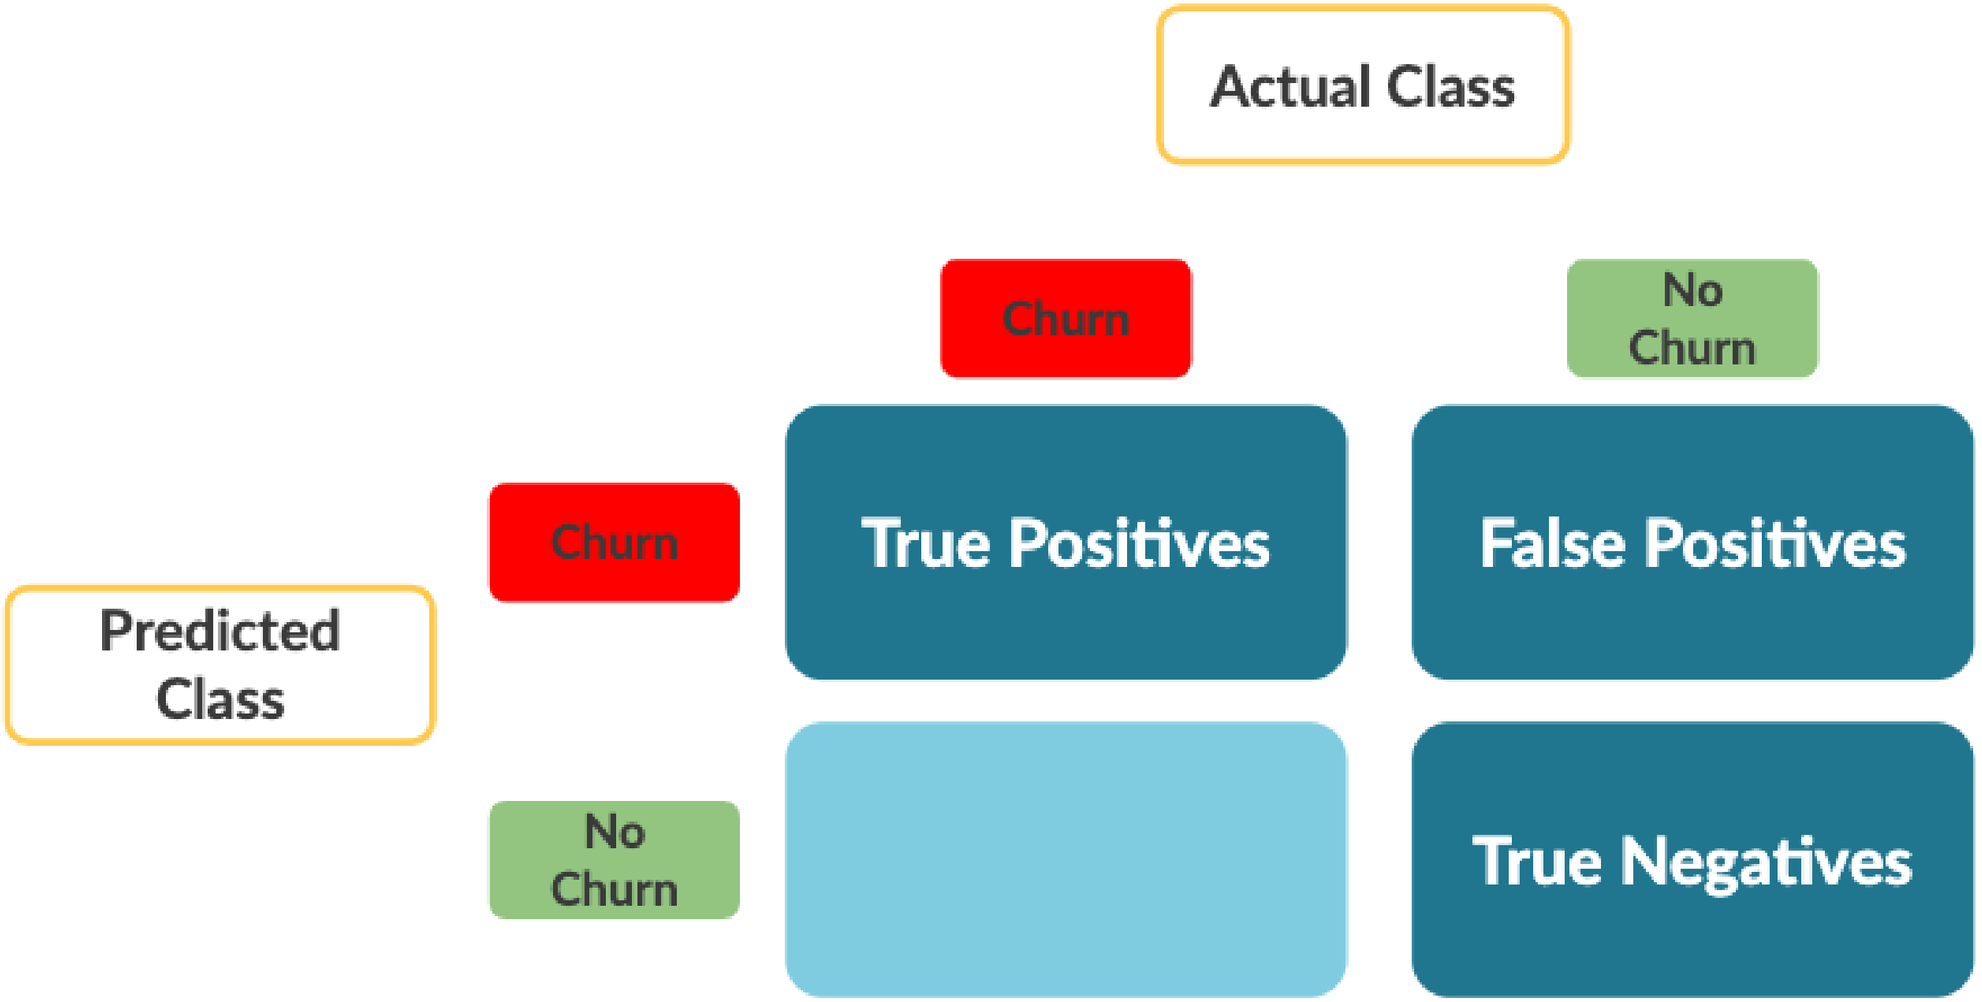
Actual Class
Churn
No Churn
Predicted Class
Churn
True Positives
False Positives
No Churn
True Negatives

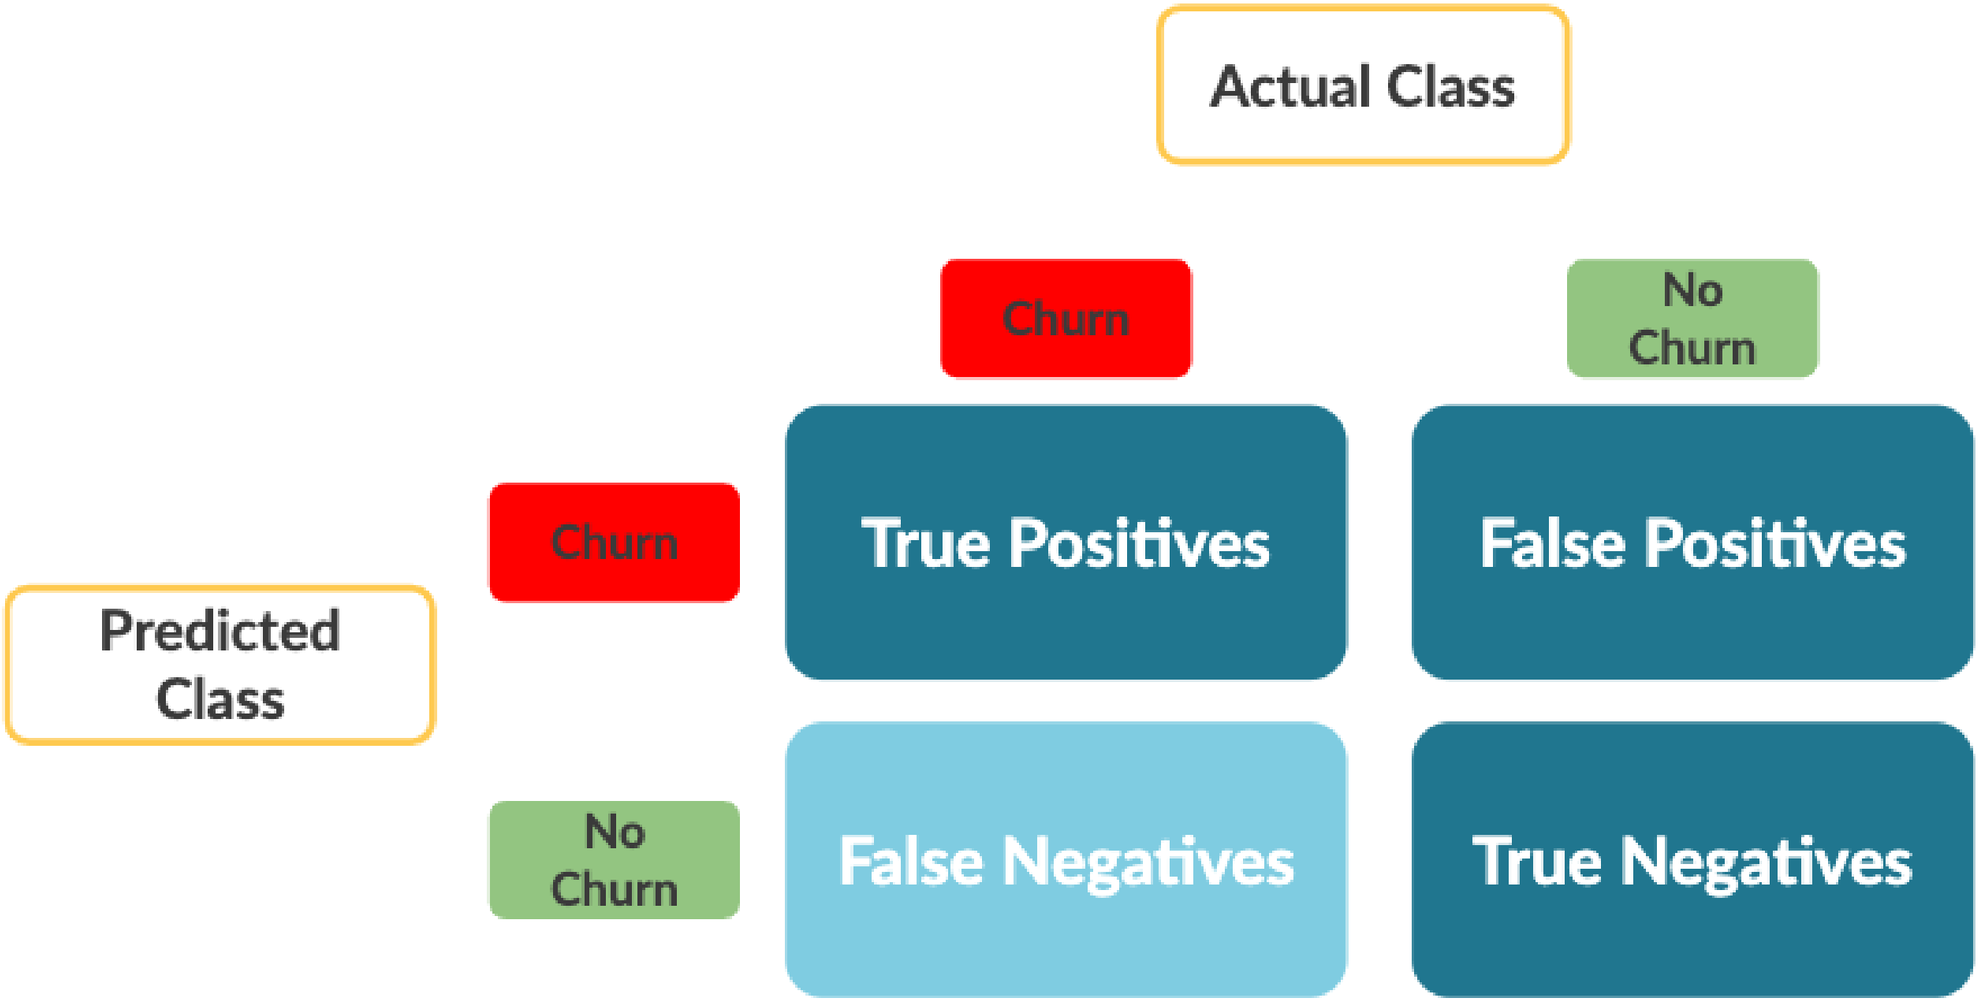
Actual Class
Churn
No Churn
Predicted Class
Churn
True Positives
False Positives
No Churn
False Negatives
True Negatives

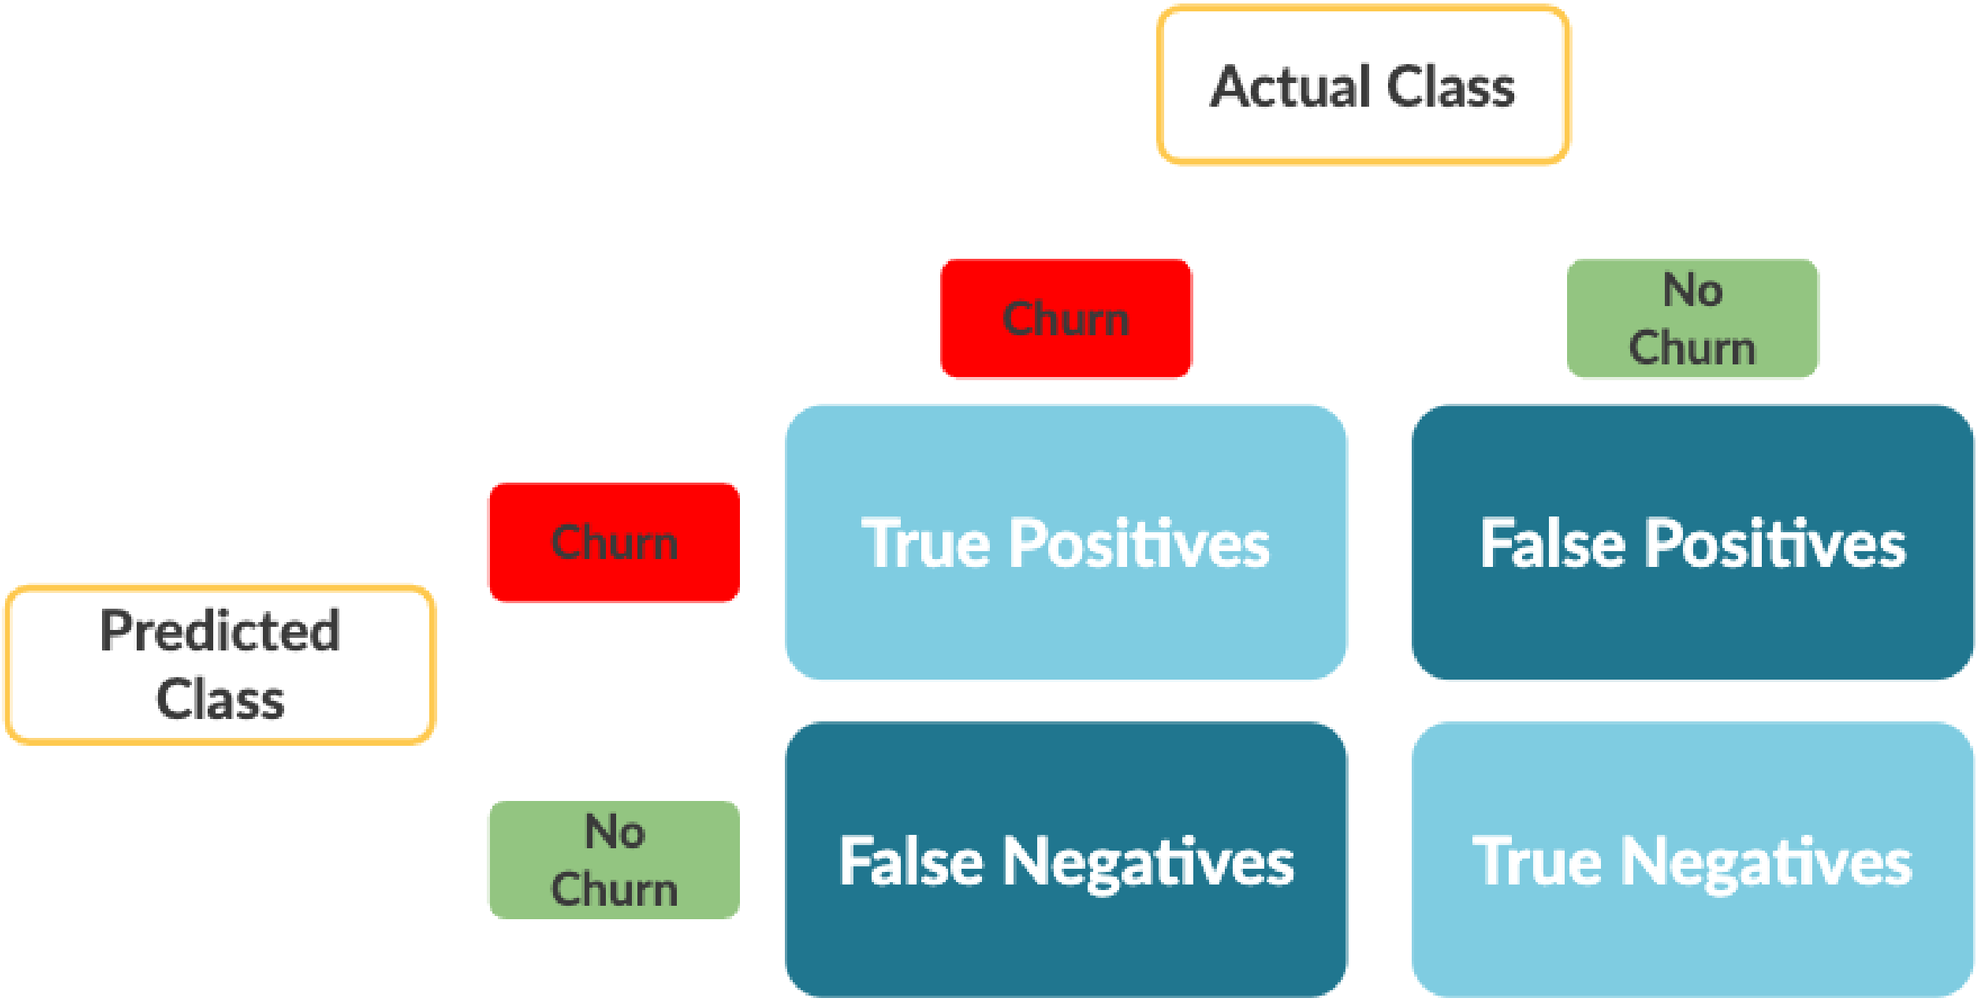
Actual Class
Churn
No Churn
Predicted Class
Churn
True Positives
False Positives
No Churn
False Negatives
True Negatives

Precision

Metric	Formula
Precision	True Positives / (True Positives + False Positives)

- A model with high precision indicates:
  - Few false positives ("false alarms")
  - Not many non-churners were classified as churners

Recall

Metric	Formula
Recall/Sensitivity	True Positives / (True Positives + False Negatives)

- A model with high recall indicates that it correctly classified most churners

Precision vs. Recall

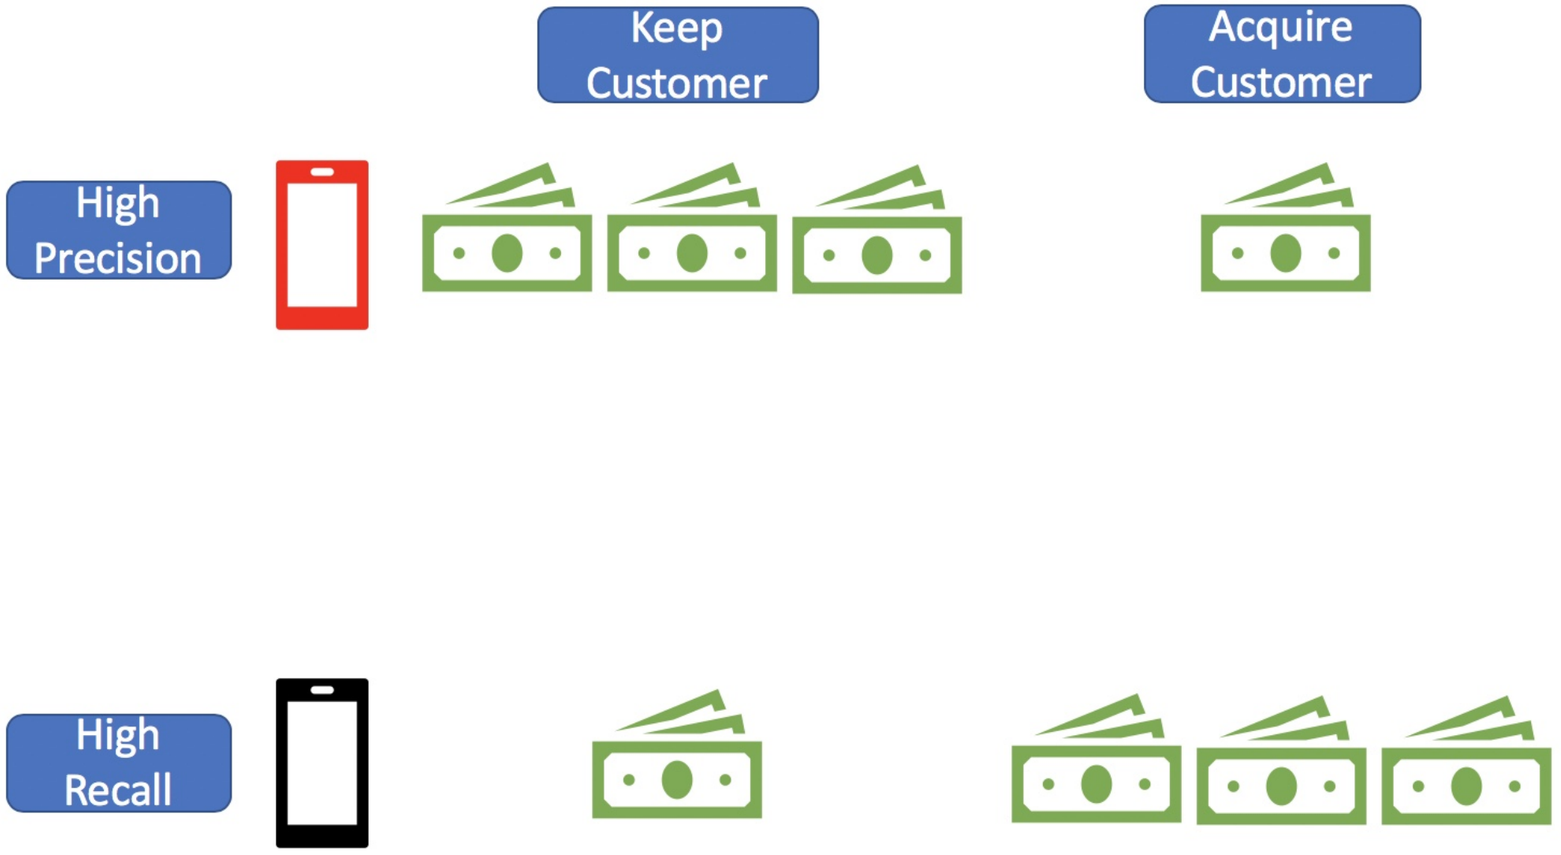

Confusion Matrix in scikit-learn

```
from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_pred)
```

Let's practice!

MARKETING ANALYTICS: PREDICTING CUSTOMER CHURN IN PYTHON

Other model metrics

MARKETING ANALYTICS: PREDICTING CUSTOMER CHURN IN PYTHON

Mark Peterson
Director of Data Science, Infoblox

Probability thresholds

- Every prediction your classifier makes has an associated probability
- Default probability threshold in scikit-learn: 50%

- What if we vary this threshold?

ROC Curve

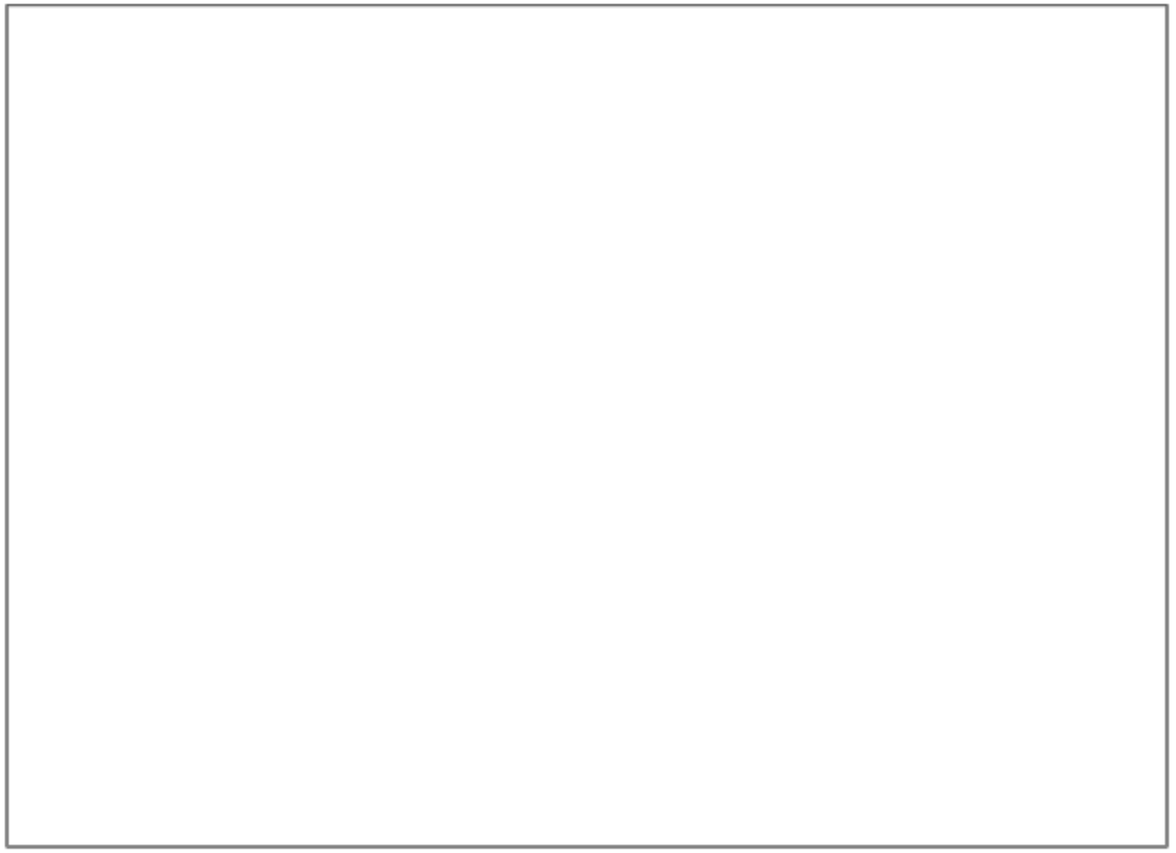

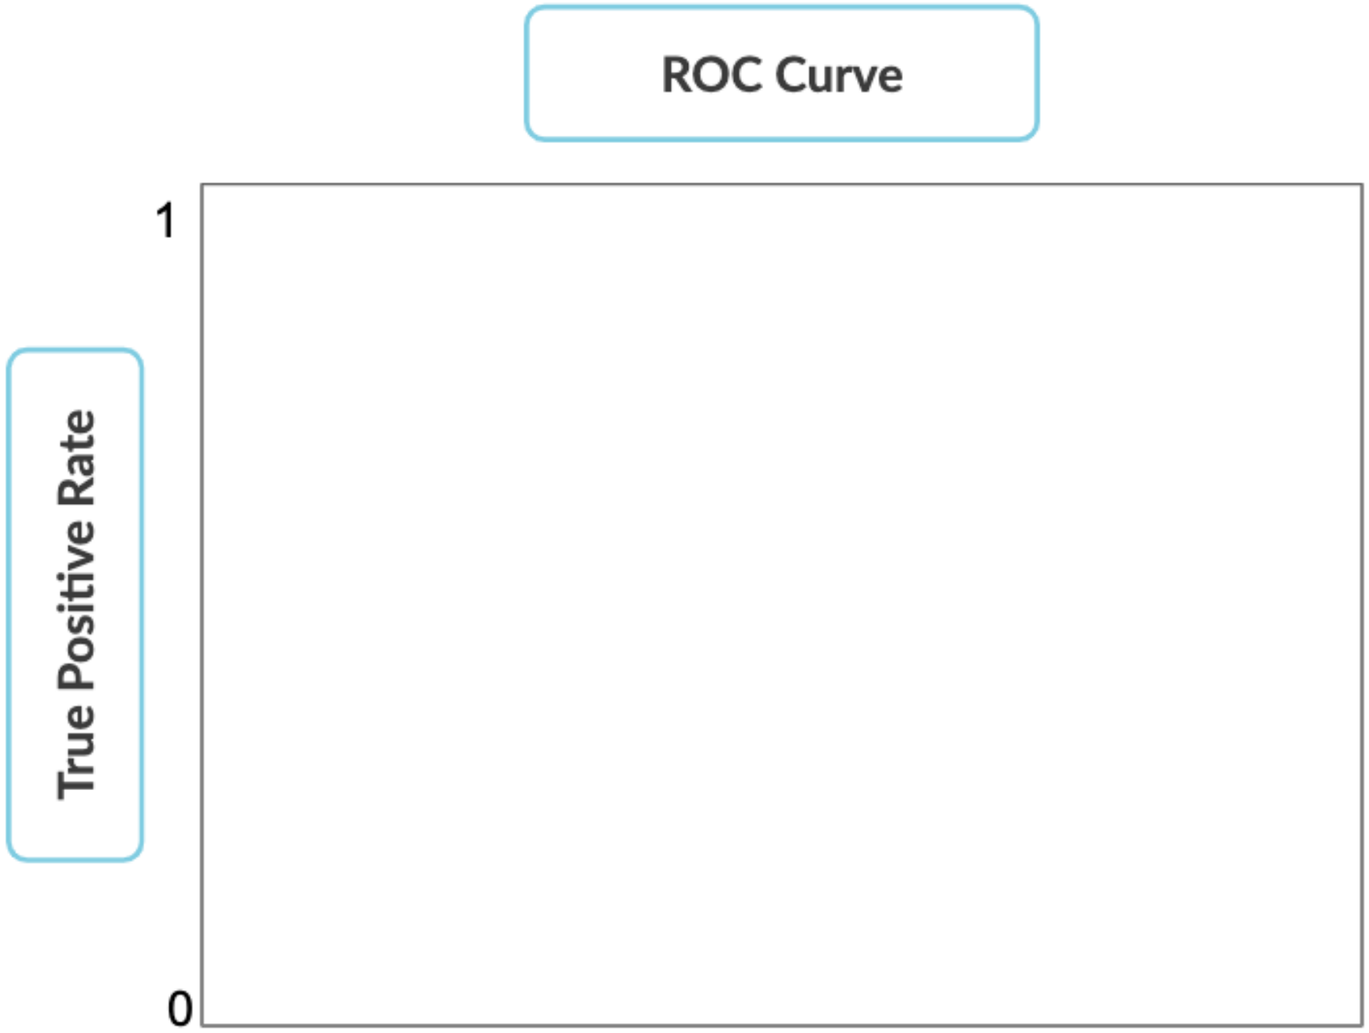
ROC Curve
True Positive Rate
1
0

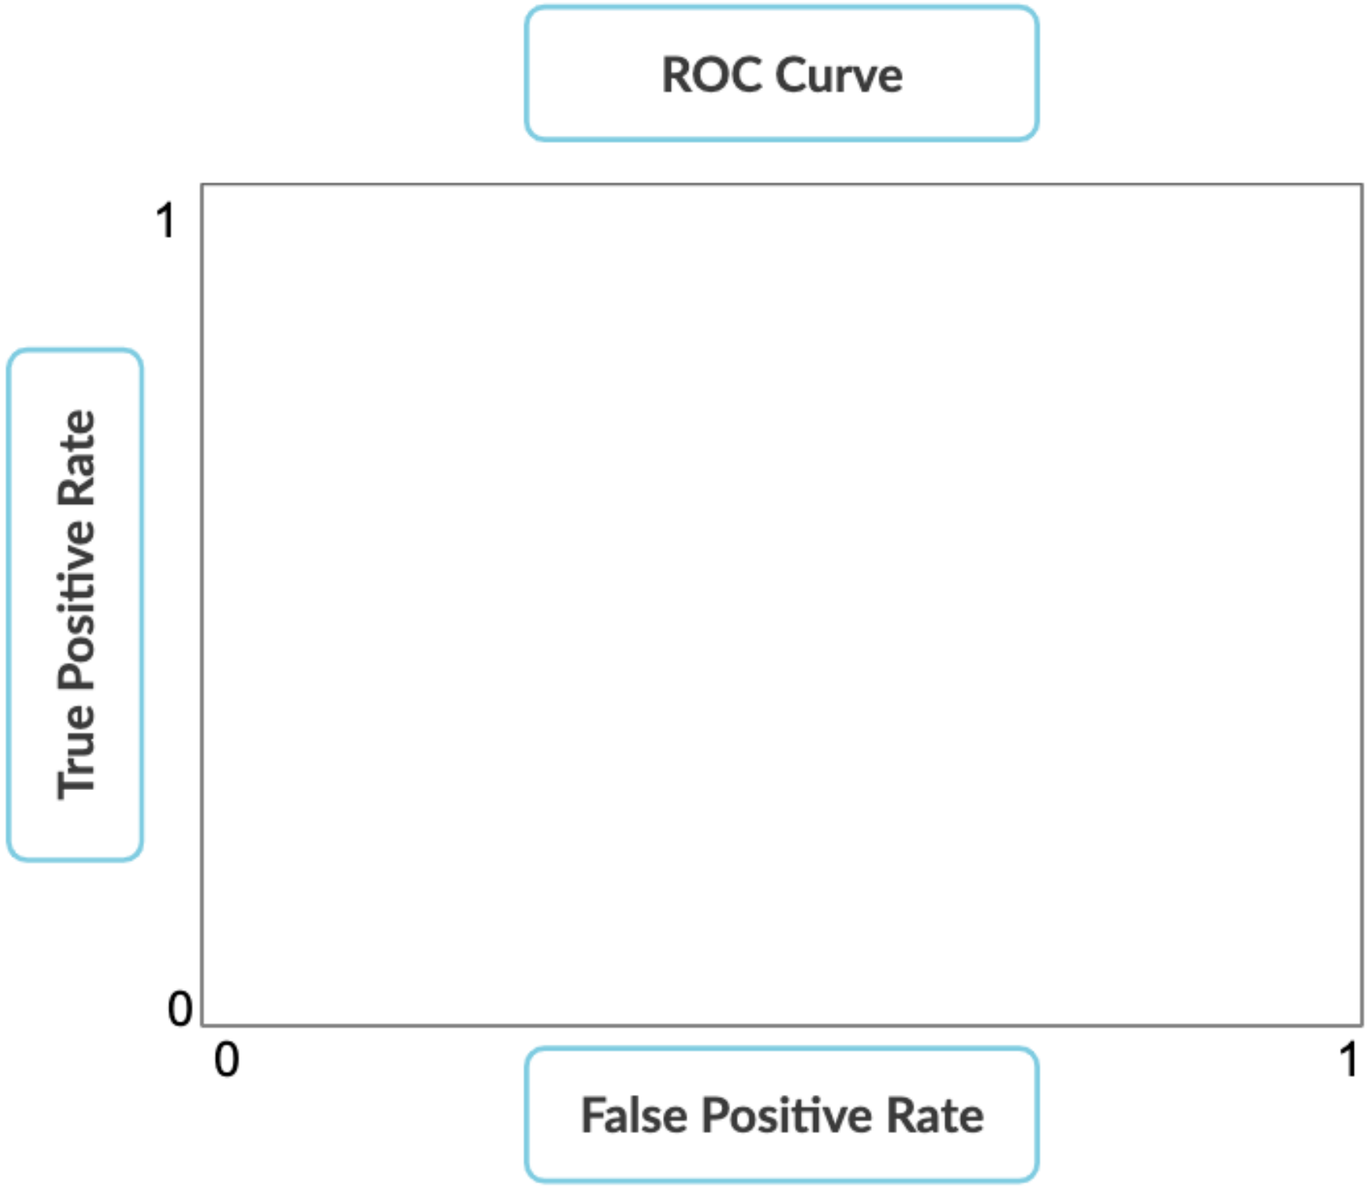
ROC Curve
True Positive Rate
1
0
0
1
False Positive Rate

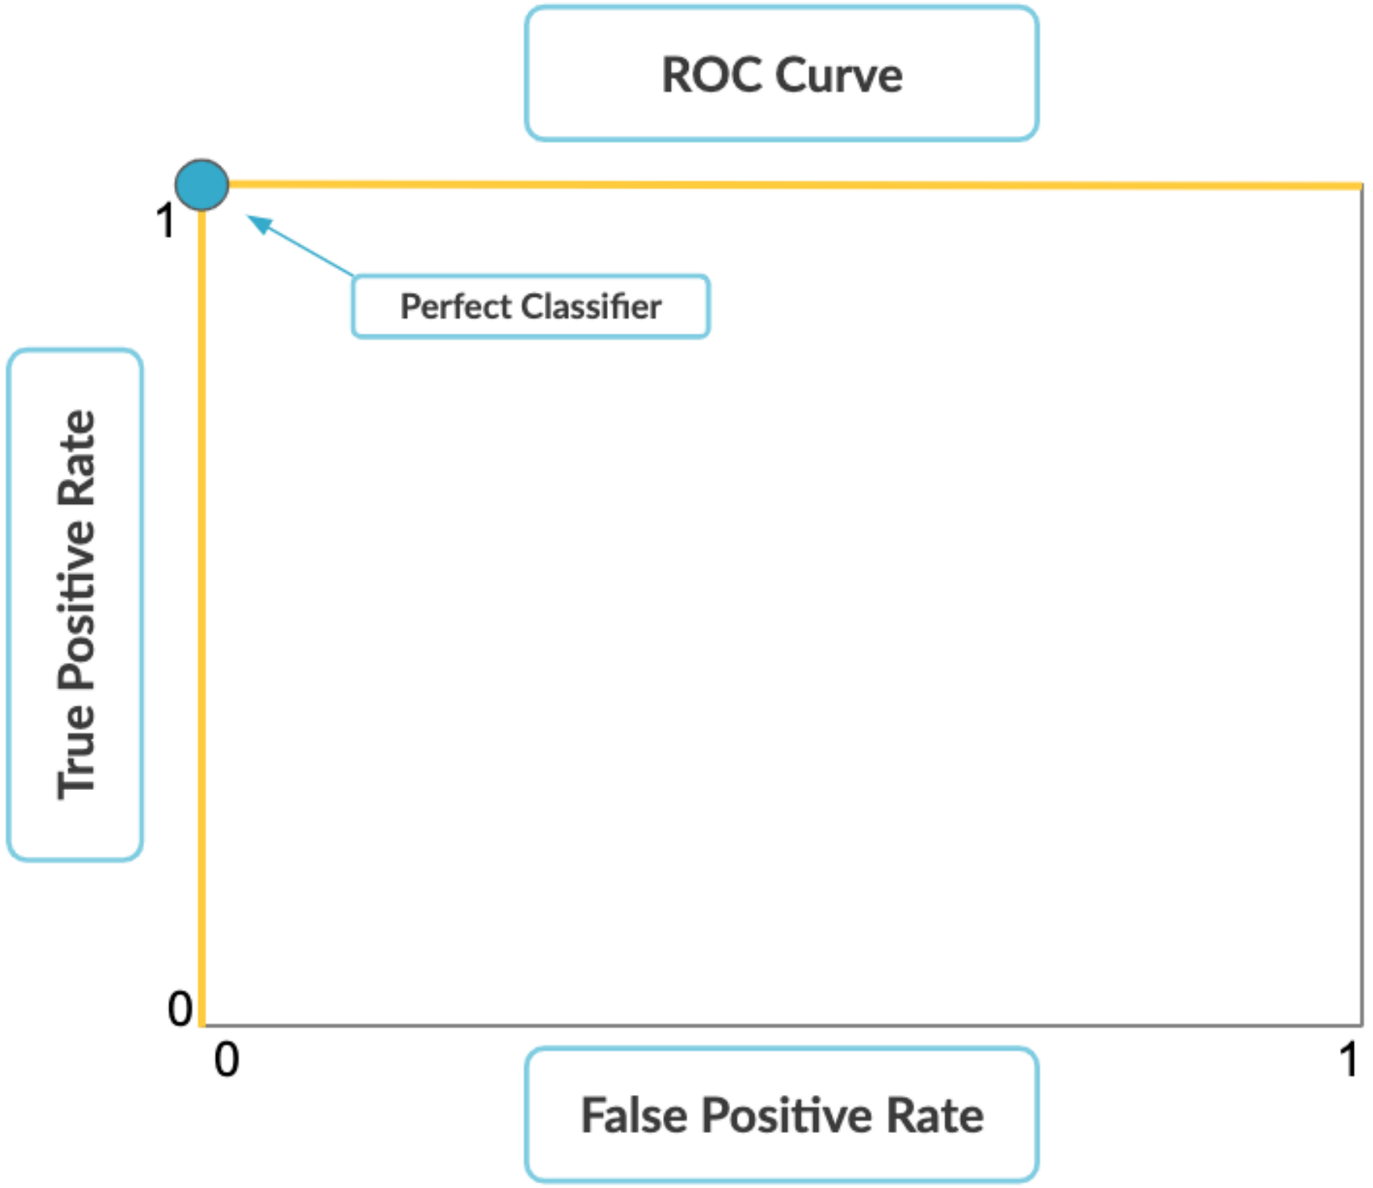
ROC Curve
1
Perfect Classifier
True Positive Rate
0
0
1
False Positive Rate

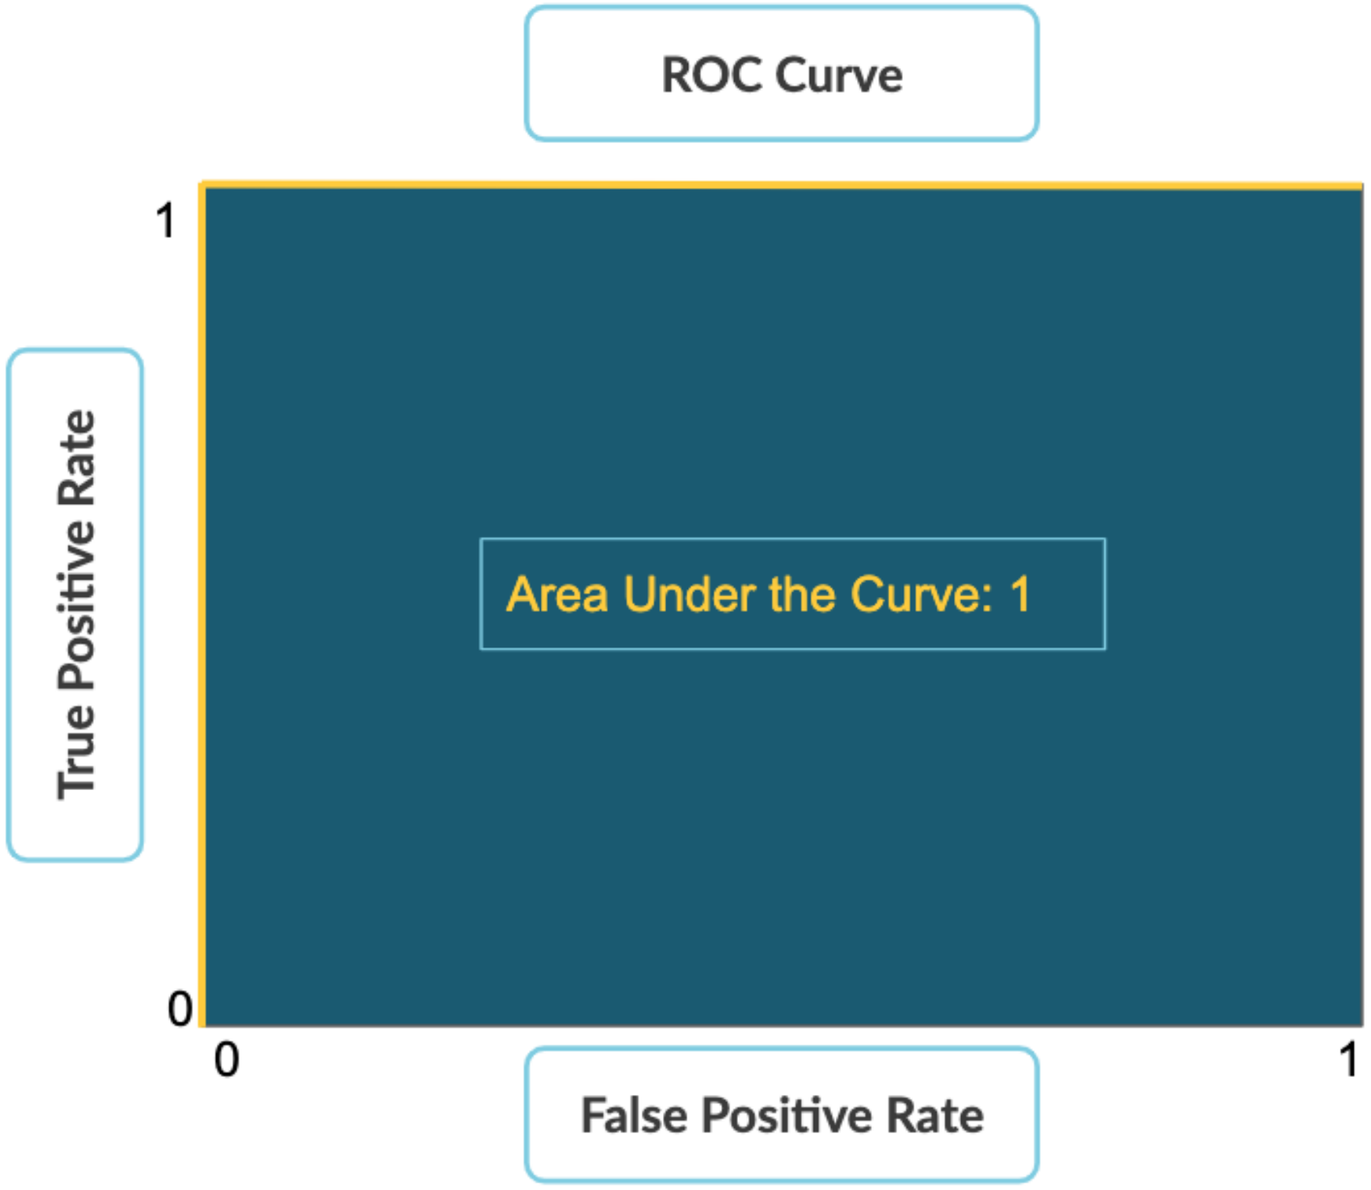
ROC Curve
1
True Positive Rate
Area Under the Curve: 1
0
0
1
False Positive Rate

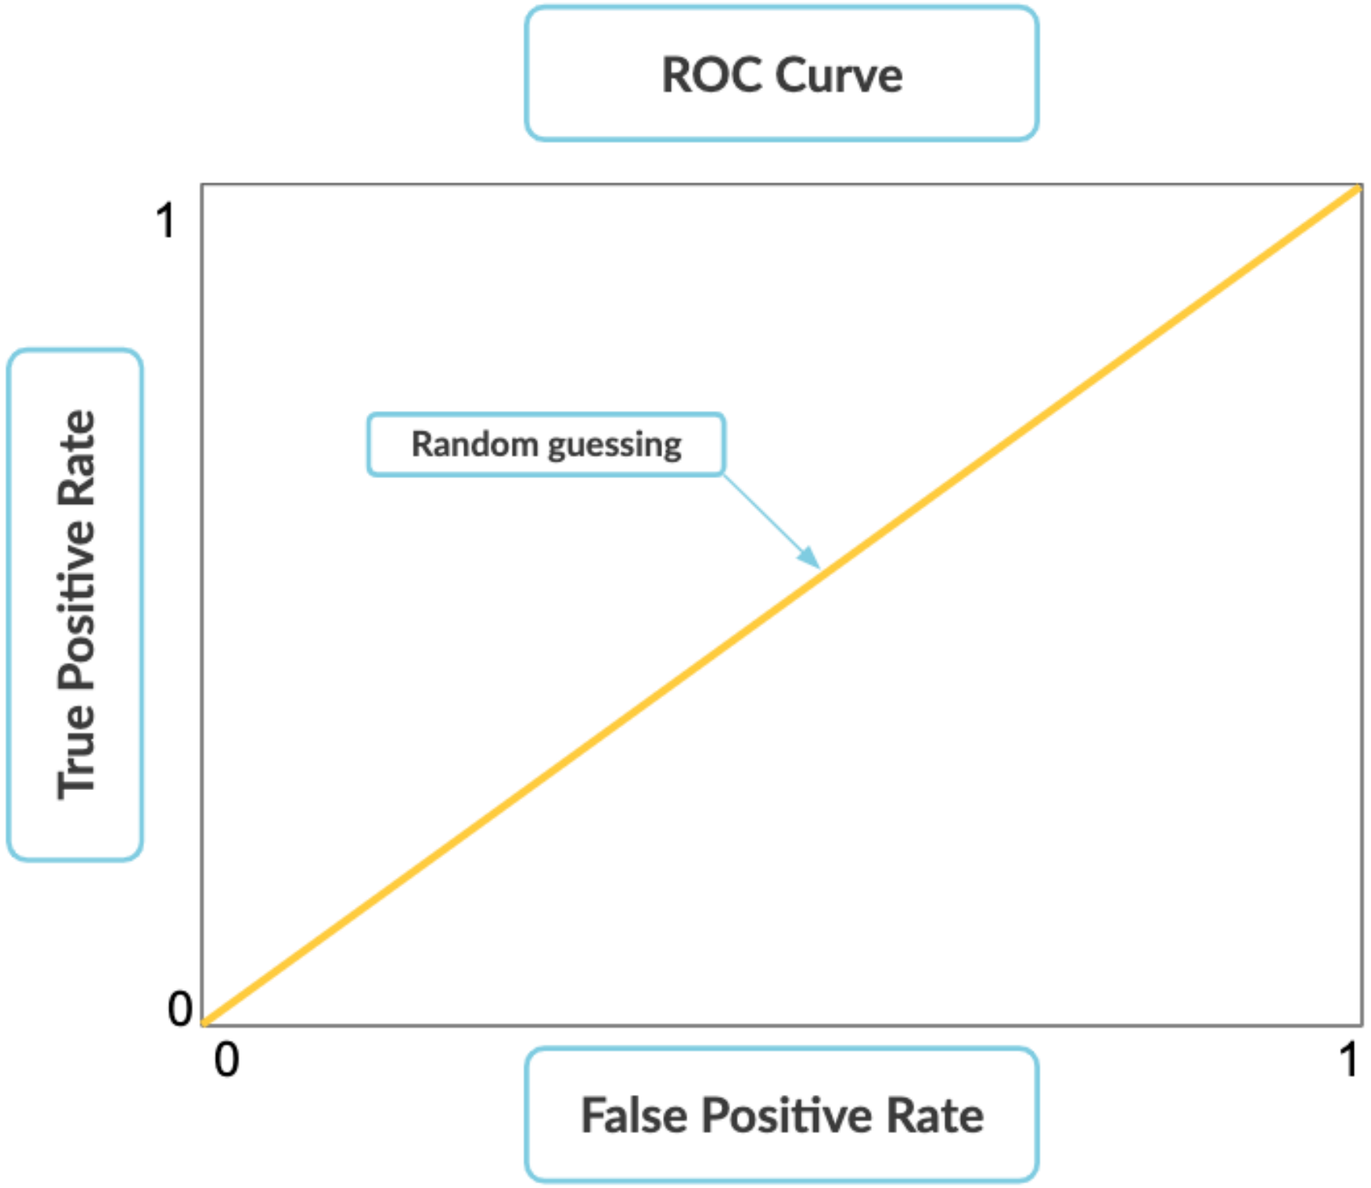
ROC Curve
1
True Positive Rate
Random guessing
0
0
1
False Positive Rate

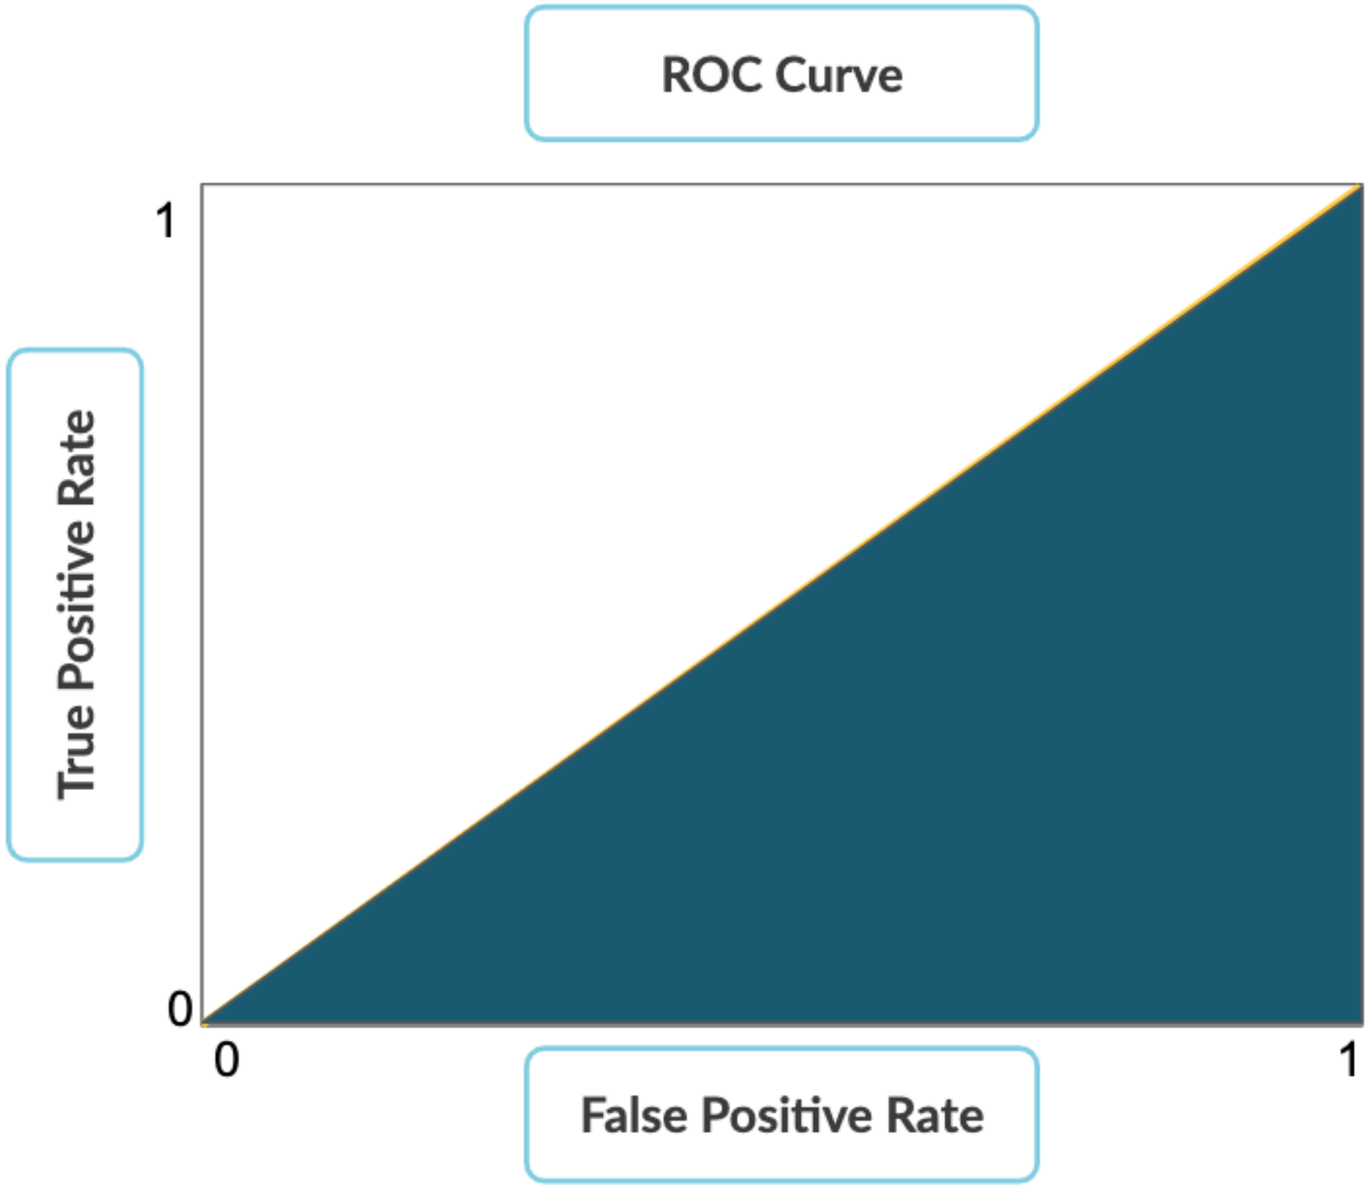
ROC Curve
True Positive Rate
1
0
0
1
False Positive Rate

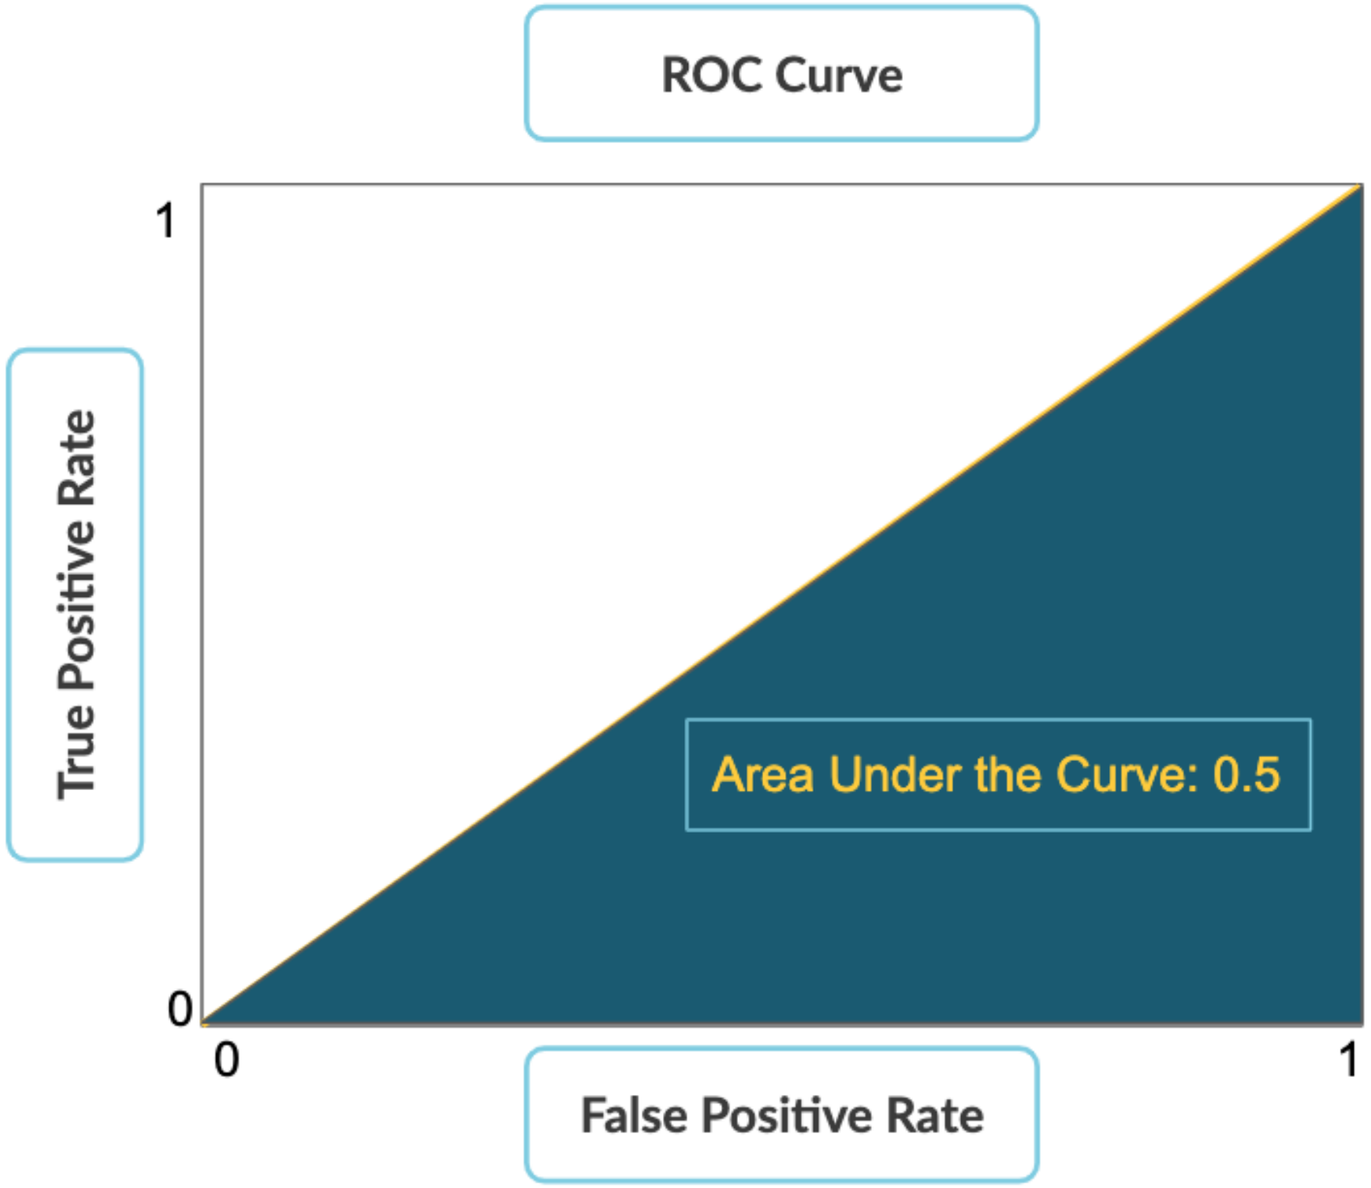
ROC Curve
1
True Positive Rate
Area Under the Curve: 0.5
0
0
1
False Positive Rate

Generating probabilities in sklearn

```
logreg.predict_proba(X_test)[:,1]
```

```
array([[0.80188981, 0.19811019],
       [0.96484075, 0.03515925],
       [0.9182671 , 0.0817329 ],
       ...,
```

```
y_pred_prob = logreg.predict_proba(X_test)[:,1]
```

ROC curve in sklearn

```
from sklearn.metrics import roc_curve

fpr, tpr, thresholds = roc_curve(y_test, y_pred_prob)
```

```
import matplotlib.pyplot as plt

plt.plot(fpr, tpr)

plt.xlabel("False Positive Rate")

plt.ylabel("True Positive Rate")

plt.plot([0, 1], [0, 1], "k--")

plt.show()
```

Area under the curve

```
from sklearn.metrics import roc_auc_score


auc = roc_auc_score(y_test, y_pred)
```

Let's practice!

MARKETING ANALYTICS: PREDICTING CUSTOMER CHURN IN PYTHON